

# CSE 1325: Object-Oriented Programming

## Lecture 17 – Chapter 10

# Multiple Inheritance, Input / Output, and Files

**Mr. George F. Rice**  
[george.rice@uta.edu](mailto:george.rice@uta.edu)

Based on material by Bjarne Stroustrup  
[www.stroustrup.com/Programming](http://www.stroustrup.com/Programming)

ERB 402  
Office Hours:  
Tuesday Thursday 11 - 12  
Or by appointment

# Quick Review

- True or **False**: A class (type) may not be used to instantiate templates such as vector
- **True** or False: The only difference in C++ between a struct and a class is the default visibility of its members.
- Ensuring that a program operates on clean, correct and useful data is **data validation**.
- Algorithmically enforceable constraints on the correctness, meaningfulness, and security of input data is **validation rules**.
- Which of the following help with that:(1) Clever programming  
**(2) Invariant classes (3) Validation methods (4) Validation in the constructor**  
(5) Limiting members to primitives
- The following are good reasons to make class fields private:  
**(1) To provide a clean interface** (2) To hide proprietary data  
**(3) To maintain an invariant** (4) To enforce HIPPA regulations  
**(5) To ease debugging**  
**(6) To change representation or algorithm without impacts**

# Quick Review

- Providing a user-defined meaning to a pre-defined operator (e.g., +, ==, <<) for a user-defined type (class) is **operator overloading**.
- Which of the following operators can NOT be overloaded: (1) :: (2) new (3) ?: (4) && **(This will NOT be on the exam!)**
- To overload the + operator for a class, define the **operator+** method.
- Reuse and extension of fields and method implementations from one class to other classes is called **inheritance**. The original class is called the **base** class, and the other classes are called **derived** classes.
- A **class hierarchy** (or **class diagram**) defines the inheritance relationships between classes.
- In specifying a base class in the derived class, the keyword “public”, “protected”, or “private” is specified. What are the implications of this?  
**“public” maintains the same visibility of inherited members.**  
**“protected” changes inherited public members to protected.**  
**“private” changes inherited public and protected members to private.**

# Overview

- Inheritance / Polymorphism Redux
- Multiple Inheritance
  - Practical Uses
  - Resolving the “diamond problem”
- Fundamental I/O concepts
  - Files and streams
  - Handling I/O errors
  - Managing stream status



# A Quick Review of Inheritance and Polymorphism

# Without Inheritance and Polymorphism

```
#include <iostream>, <vector>, <chrono>, <thread>
using namespace std;

class Critter {
public:
    Critter(int frequency) : _frequency{frequency}, _timer{frequency} { }
    void count() {if (++_timer > _frequency) _timer = 0;}
    void speak() { if (!_timer) cout << "Generic critter sound!" << endl; }
protected:
    int _frequency;
    int _timer;
};
int main() {
    vector<Critter*> critters;

    critters.push_back(new Critter{13});
    critters.push_back(new Critter{11});
    critters.push_back(new Critter{7});
    critters.push_back(new Critter{3});

    cout << "W E L C O M E   T O   T H E   B A R N Y A R D !" << endl;
    for (int i=0; i<120; ++i) {
        for (Critter *c: critters) {
            c->count();
            c->speak();
        }
        this_thread::sleep_for(chrono::milliseconds(50));
    }
}
```

What code needs to change  
to add barnyard animals  
(cow, chicken, dog...)?

# With Inheritance and Polymorphism

```
#include <iostream>, <vector>, <chrono>, <thread>
using namespace std;

class Critter {
public:
    Critter(int frequency) : _frequency{frequency}, _timer{frequency} { }
    void count() {if (++_timer > _frequency) _timer = 0;}
    virtual void speak() { if (!_timer) cout << "Generic critter sound!" << endl; }
protected:
    int _frequency;
    int _timer;
};

class Cow : public Critter {
public:
    Cow(int frequency) : Critter(frequency) { }
    void speak() override { if (!_timer) cout << "Moo! Mooooo!" << endl; }
};

class Dog : public Critter {
public:
    Dog(int frequency) : Critter(frequency) { }
    void speak() override { if (!_timer) cout << "Woof! Woof!" << endl; }
};

class Chicken : public Critter {
public:
    Chicken(int frequency) : Critter(frequency) { }
    void speak() override { if (!_timer) cout << "Cluck! Cluck!" << endl; }
};
```

What code needs to change  
to add barnyard animals  
(cow, chicken, dog...)?

# With Inheritance and Polymorphism

```
int main() {
    vector<Critter*> critters;

    critters.push_back(new Critter{13});
    critters.push_back(new Cow{11});
    critters.push_back(new Dog{7});
    critters.push_back(new Chicken{3});

    cout << "W E L C O M E   T O   T H E   B A R N Y A R D !" << endl;
    for (int i=0; i<120; ++i) {
        for (Critter *c: critters) {
            c->count();
            c->speak();
        }
        this_thread::sleep_for(chrono::milliseconds(50));
    }
}
```

Literally 4 words changed!

Added 3 derived classes – but mostly just the portions that differed from the base class.

What code needs to change to add barnyard animals (cow, chicken, dog...)?

```
ricegf@pluto:~/dev/cpp/201701/17$ make v2
g++ -std=c++11 -o v2 critter_v2.cpp
ricegf@pluto:~/dev/cpp/201701/17$ ./v2
W E L C O M E   T O   T H E   B A R N Y A R D !
Generic critter sound!
Moo! Mooooo!
Woof! Woof!
Cluck! Cluck!
Cluck! Cluck!
Woof! Woof!
Cluck! Cluck!
Moo! Mooooo!
Cluck! Cluck!
Generic critter sound!
Woof! Woof!
Cluck! Cluck!
```

# Multiple Inheritance

- What if a class is derived from more than one base class?
  - This is called *multiple inheritance*
  - MI is often helpful, e.g., with the Adapter Pattern
    - Given the Client and Server classes, simply inherit both!
- With multiple inheritance, each base class's members are laid out sequentially in memory after the derived class's members

**Multiple Inheritance** is a derived class inheriting class members from two or more base classes.



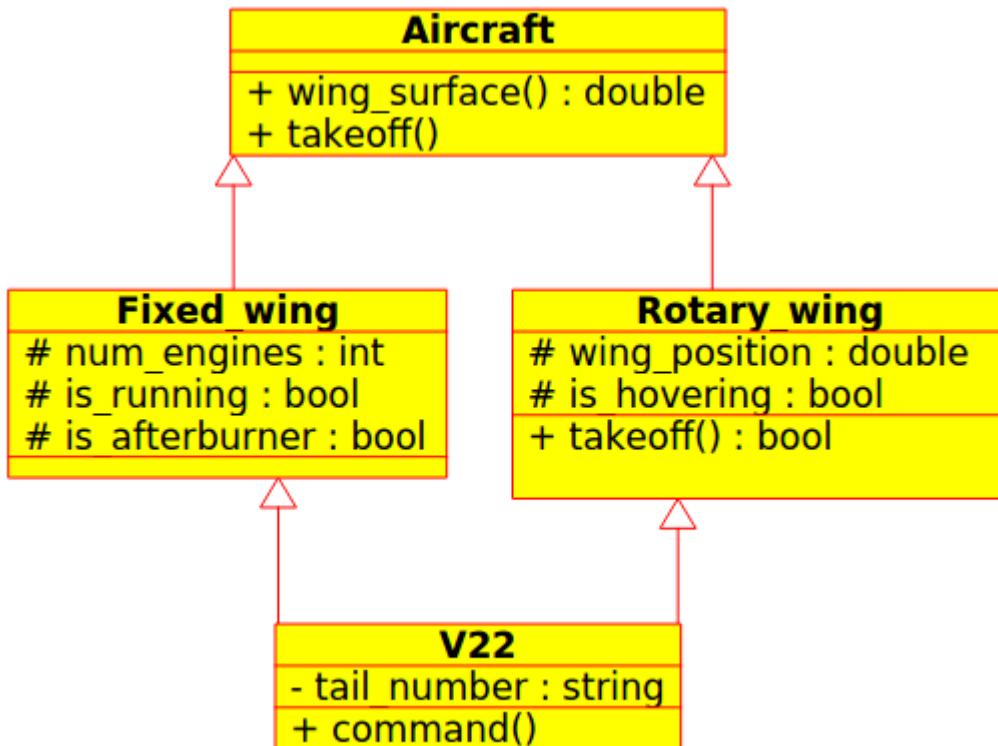
# Multiple Inheritance

V-22 Tilt Rotor Aircraft



Image (c) 2006, Greg Heartsfield – used under the GNU Free Documentation License  
[https://commons.wikimedia.org/wiki/File:V-22\\_preparation\\_alliance\\_airport.jpg](https://commons.wikimedia.org/wiki/File:V-22_preparation_alliance_airport.jpg)

# Example with Multiple Inheritance

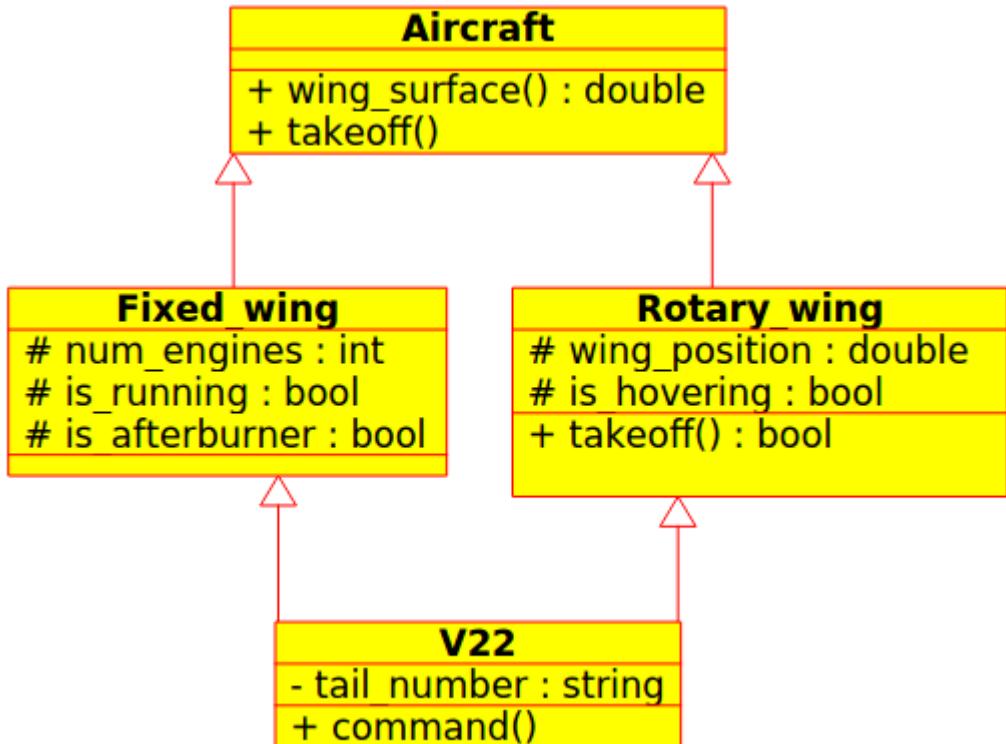


V22 would have `wing_surface()`, `takeoff()`, and `command()` methods

**Wait a minute...  
which `takeoff()`?**

And, more subtly, both the **Fixed\_wing** and **Rotary\_wing** classes inherit `wing_surface()`. Which path will inheritance follow, i.e., *which* inherited `wing_surface()` will V22 invoke?

# Example with Multiple Inheritance



This is called the “**diamond problem**” - when inheriting in a “diamond shape”, inheritance paths are no longer obvious.

Java, C#, and Ruby (for example) sidestep this problem and do **NOT** implement multiple inheritance, using Interfaces (also called Protocols) instead.

Python makes inheritance order significant, and calls the first method found – left to right, then bottom to top.

How does C++ react?

# The ABCD Diamond

```
#include <iostream>
using namespace std;

class A {
public:
    virtual void m() {cout << "m of A" << endl;}
};

class B : public A {
public:
    virtual void m() override {cout << "m of B" << endl;}
};

class C : public A { };

class D : public B, public C { };

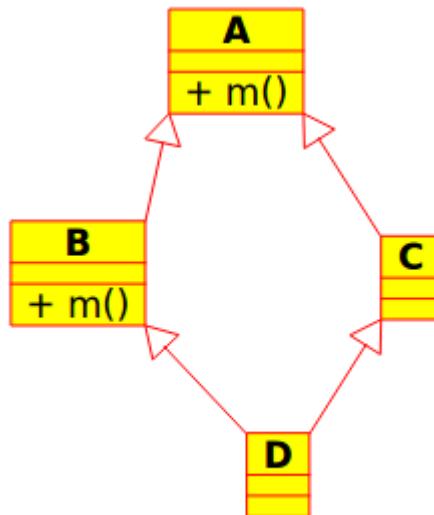
int main() {
    D d;
    d.m();
}
```

C++ (for once) refuses to make an assumption

```
ricegf@pluto:~/dev/cpp
diamond.cpp: In function
diamond.cpp:22:5: error
    d.m();
```

C++ (for once) refuses to make an assumption and instead raises a compiler error.

Let's simplify and test  
(a GREAT strategy for  
learning a new language!).



# Explicit Base Class Method Call

```
#include <iostream>
using namespace std;

class A {
public:
    virtual void m() {cout << "m of A" << endl;}
};

class B : public A {
public:
    virtual void m() override {cout << "m of B" << endl;}
};

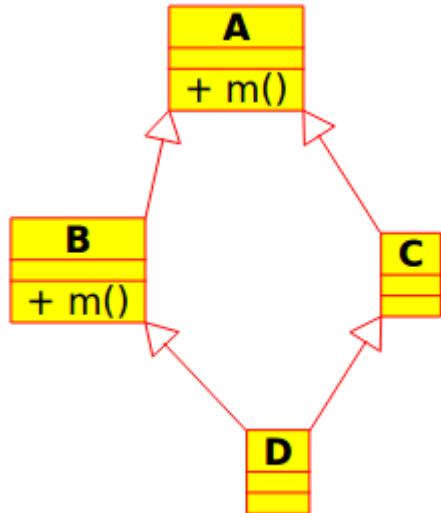
class C : public A { };

class D : public B, public C { };

int main() {
    D d;
    d.B::m();
    d.C::m();
}
```

Ambiguity resolved!

What if we specify the method using namespace notation?



```
ricegef@pluto:~/dev/cpp/201701/17$ g++ -std=c++11 diamond.cpp
ricegef@pluto:~/dev/cpp/201701/17$ ./a.out
m of B
m of A
ricegef@pluto:~/dev/cpp/201701/17$ █
```

# The ABCD Diamond

```
#include <iostream>
using namespace std;

class A {
public:
    virtual void m() {cout << "m of A" << endl;}
};

class B : public A {
public:
    virtual void m() override {cout << "m of B" << endl;}
};

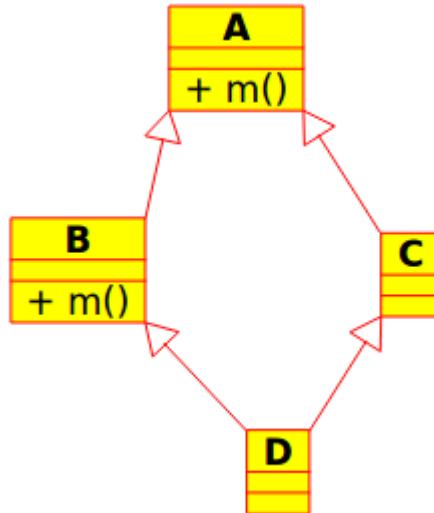
class C : public A { };

class D : public B, public C { };

int main() {
    D d;
    d.A::m();
    d.B::m();
    d.C::m();
}
```

Uh oh

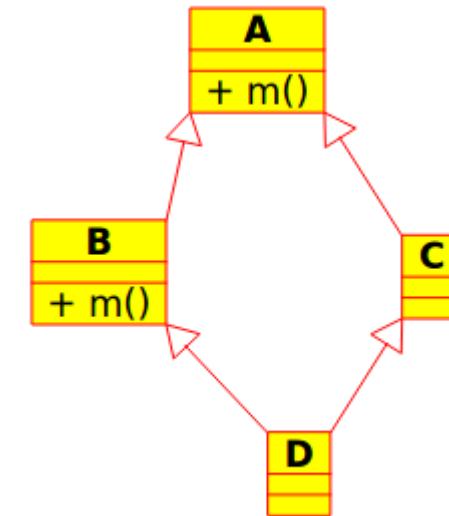
Let's explicitly call A's method, then...



```
ricegf@pluto:~/dev/cpp/201701/17$ g++ -std=c++11 diamond.cpp
diamond.cpp: In function ‘int main()’:
diamond.cpp:20:8: error: ‘A’ is an ambiguous base of ‘D’
      d.A::m();
            ^
ricegf@pluto:~/dev/cpp/201701/17$
```

# This is a Memory Layout Issue

- B adds A's method to the end of its memory block
- C adds A's method to the end of its memory block
- D now has two A's – the one in B and the one in C – and doesn't know which A you mean
- Proposed solution: Use pointers and only keep a single A!



Memory Layout

Memory for D  
Memory for B  
method m  
**Memory for A**  
method m  
Memory for C  
method m  
**Memory for A**  
method m

# Virtual Inheritance

```
#include <iostream>
using namespace std;

class A {
public:
    virtual void m() {cout << "m of A" << endl;}
};

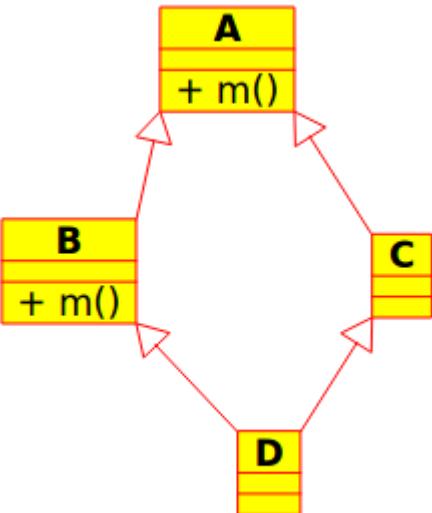
class B : public virtual A {
public:
    virtual void m() override {cout << "m of B" << endl;}
};

class C : public virtual A { };

class D : public B, public C { };

int main() {
    D d;
    d.A::m();
    d.B::m();
    d.C::m();
}
```

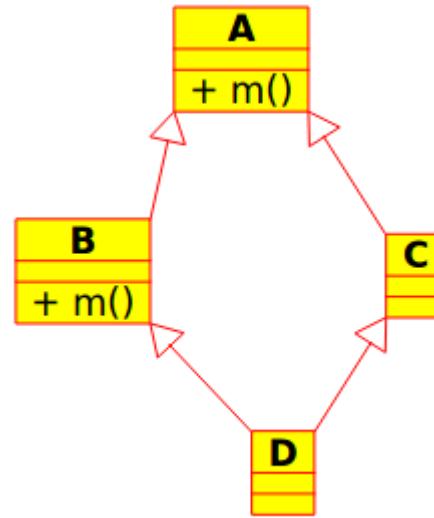
Easier done than said.  
Just make the intermediate  
classes inherit *virtually*.



```
ricegf@pluto:~/dev/cpp/201701/17$ g++ -std=c++11 diamond.cpp
ricegf@pluto:~/dev/cpp/201701/17$ ./a.out
m of A
m of B
m of A
ricegf@pluto:~/dev/cpp/201701/17$ █
```

# How Does Virtual Inheritance Work?

- By keeping a single block of memory allocated for A, the ambiguity is resolved



Memory for D  
Memory for B  
method m  
pointer to A  
Memory for C  
method m  
pointer to A

Memory for A  
method m

# Object Layout in Memory

```
class Parent1 {  
public:  
    int a1, b1;  
    virtual void foo()  
        {cout << "parent 1" << endl;}  
};  
class Parent2 {  
public:  
    int a2, b2;  
    virtual void foo()  
        {cout << "parent 2" << endl;}  
};  
class Child : public Parent1,  
              public Parent2 {  
public:  
    int c, d;  
    virtual void foo()  
        {cout << "child" << endl;}  
};  
int main(){  
    Parent1 p1;  
    Parent2 p2;  
    Child c;  
  
    p1.foo();  
    p2.foo();  
    c.foo();
```

```
cout << "Parent1 Offset a1 = "  
     << (size_t)&p1.a1 - (size_t)&p1 << endl;  
cout << "Parent1 Offset b1 = "  
     << (size_t)&p1.b1 - (size_t)&p1 << endl;  
  
cout << "Parent2 Offset a2 = "  
     << (size_t)&p2.a2 - (size_t)&p2 << endl;  
cout << "Parent2 Offset b2 = "  
     << (size_t)&p2.b2 - (size_t)&p2 << endl;  
  
cout << "Child Offset a1 = "  
     << (size_t)&c.a1 - (size_t)&c << endl;  
cout << "Child Offset b1 = "  
     << (size_t)&c.b1 - (size_t)&c << endl;  
cout << "Child Offset a2 = "  
     << (size_t)&c.a2 - (size_t)&c << endl;  
cout << "Child Offset b2 = "  
     << (size_t)&c.b2 - (size_t)&c << endl;  
cout << "Child Offset c = "  
     << (size_t)&c.c - (size_t)&c << endl;  
cout << "Child Offset d = "  
     << (size_t)&c.d - (size_t)&c << endl;
```

}

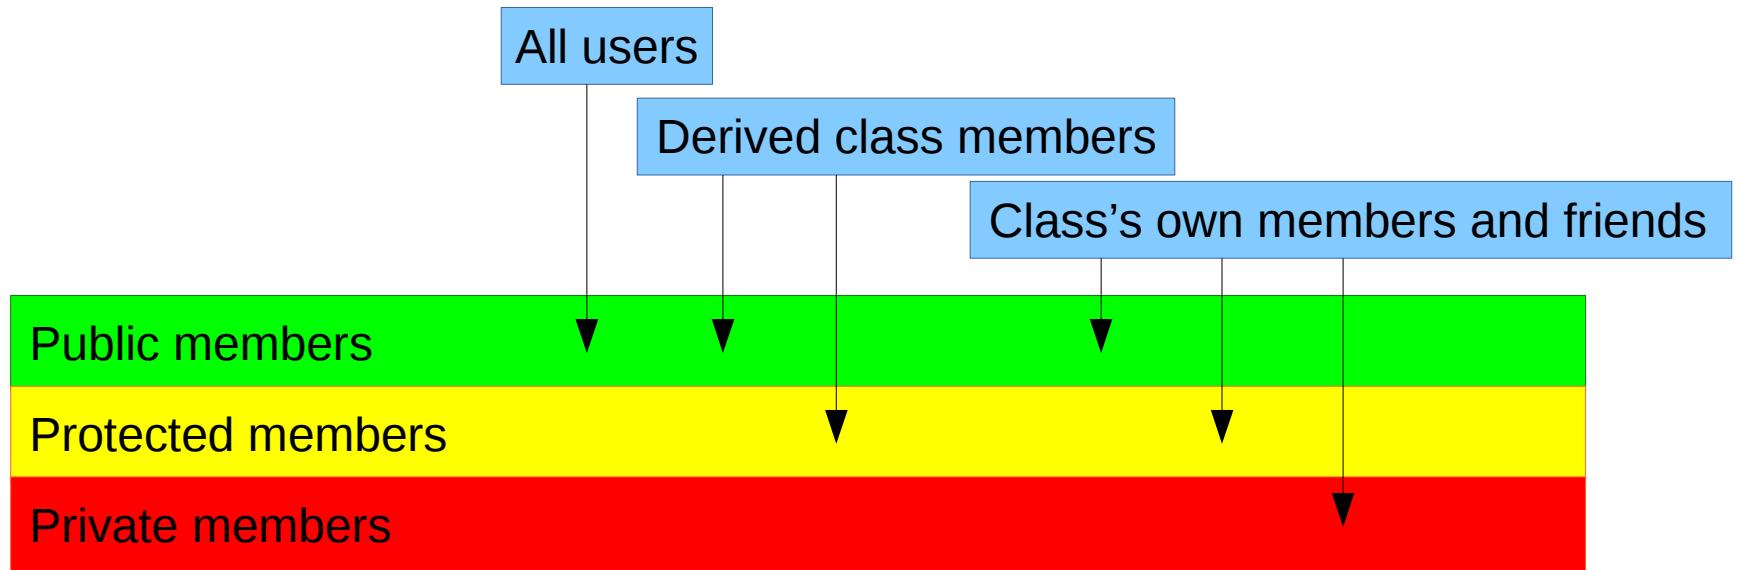
# Object Layout in Memory

```
class Parent1 {  
public:  
    int a1, b1;  
    virtual void foo()  
        {cout << "parent 1" << endl;}  
};  
class Parent2 {  
public:  
    int a2, b2;  
    virtual void foo()  
        {cout << "parent 2" << endl;}  
};  
class Child : public Parent1,  
public:  
    int c,  
    virtual child  
        {cout Parent1 Offset a1 = 8  
        Parent1 Offset b1 = 12  
};  
int main(){ Parent2 Offset a2 = 8  
Parent1 Parent2 Offset b2 = 12  
Parent2 Child Offset a1 = 8  
Child c Child Offset b1 = 12  
        Child Offset a2 = 24  
p1.foo() Child Offset b2 = 28  
p2.foo() Child Offset c = 32  
c.foo() Child Offset d = 36  
ricecgf@pluto:~/dev/cpp/19$ ./a.out
```

```
cout << "Parent1 Offset a1 = "  
    << (size_t)&p1.a1 - (size_t)&p1 << endl;  
cout << "Parent1 Offset b1 = "  
    << (size_t)&p1.b1 - (size_t)&p1 << endl;  
  
cout << "Parent2 Offset a2 = "  
    << (size_t)&p2.a2 - (size_t)&p2 << endl;  
cout << "Parent2 Offset b2 = "  
    << (size_t)&p2.b2 - (size_t)&p2 << endl;  
  
cout << "Child Offset a1 = "  
    << (size_t)&c.a1 - (size_t)&c << endl;  
cout << "Child Offset b1 = "  
    << (size_t)&c.b1 - (size_t)&c << endl;  
cout << "Child Offset a2 = "  
    << (size_t)&c.a2 - (size_t)&c << endl;  
cout << "Child Offset b2 = "  
    << (size_t)&c.b2 - (size_t)&c << endl;  
cout << "Child Offset c = "  
    << (size_t)&c.c - (size_t)&c << endl;  
cout << "Child Offset d = "  
    << (size_t)&c.d - (size_t)&c << endl;
```

**BUT**  
**Memory layout of objects  
is not specified  
in the standard,  
but is *compiler-specific!***

# Access model



- A member (data, function, or type member) or a base can be
  - Private, protected, or public

# Pure virtual functions

- Often, a function in an interface can't be implemented
  - E.g. the data needed is "hidden" in the derived class
  - We must ensure that a derived class implements that function
  - Make it a "pure virtual function" (`=0`)
- This is how we define truly abstract interfaces ("pure interfaces")

```
class Engine {      // interface to electric motors
    // no data
    // (usually) no constructor
    virtual double increase(int i) =0;      // must be defined in a derived class
    // ...
    virtual ~Engine();      // (usually) a virtual destructor
};

Engine eee;      // error: Collection is an abstract class
```

# Pure virtual functions

- A pure interface can then be used as a base class

```
Class M123 : public Engine {      // engine model M123
    // representation
public:
    M123(); // constructor: initialization, acquire resources
    double increase(int i) { /* ... */ } // overrides Engine ::increase
    // ...
    ~M123(); // destructor: cleanup, release resources
};

M123 window3_control; // OK
```



# Switching Gears Now To I/O and Files

# Input and Output

**data source:**



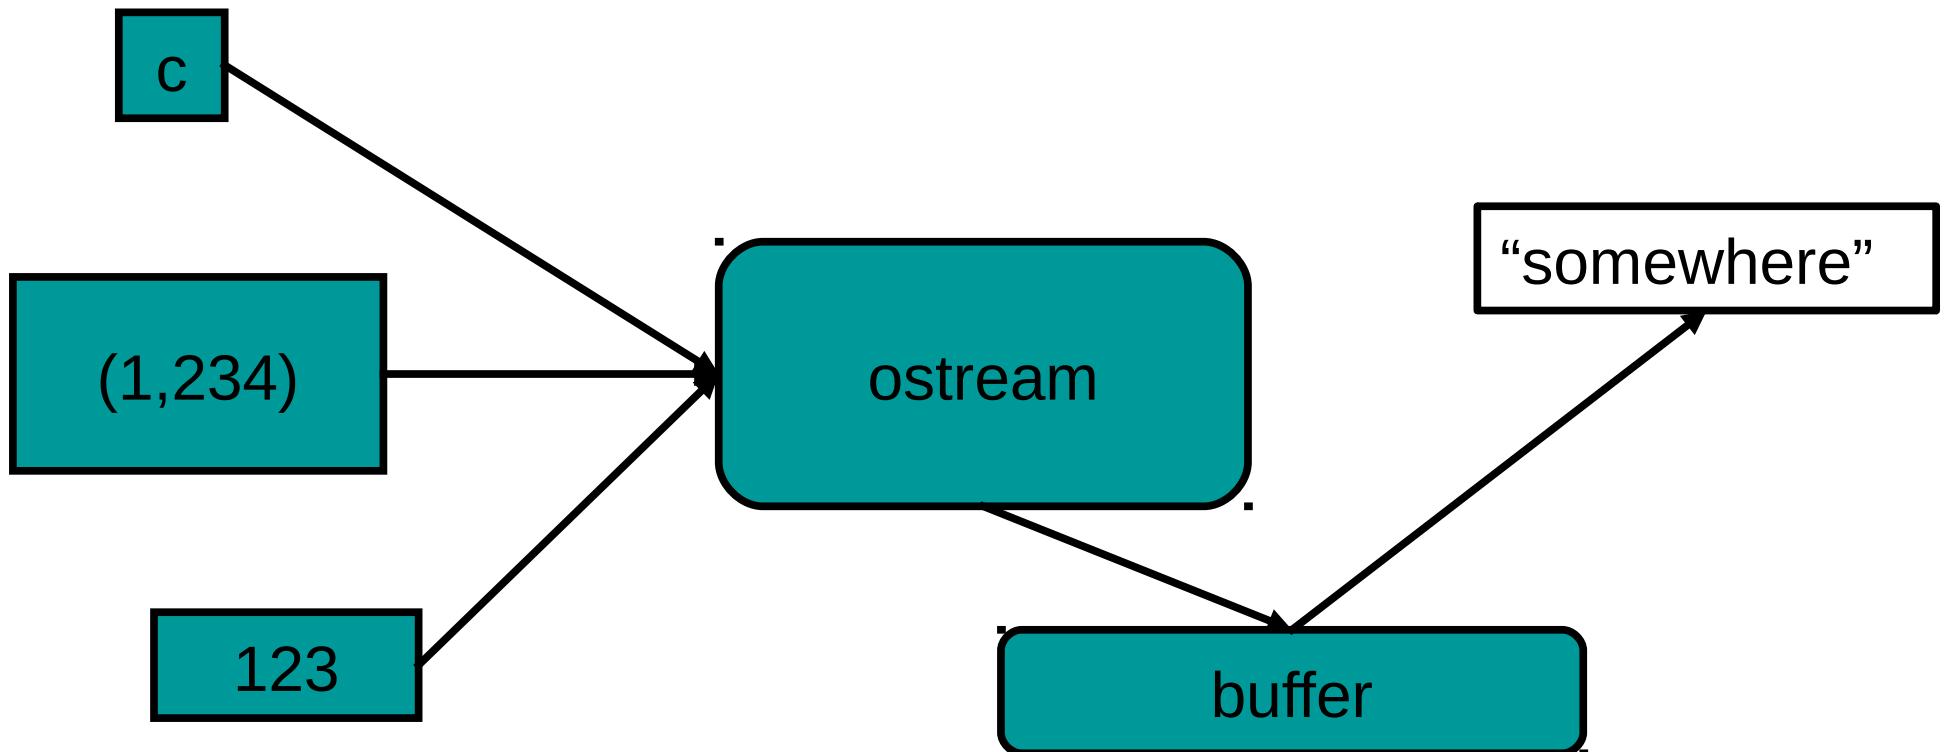
our program

**data destination:**



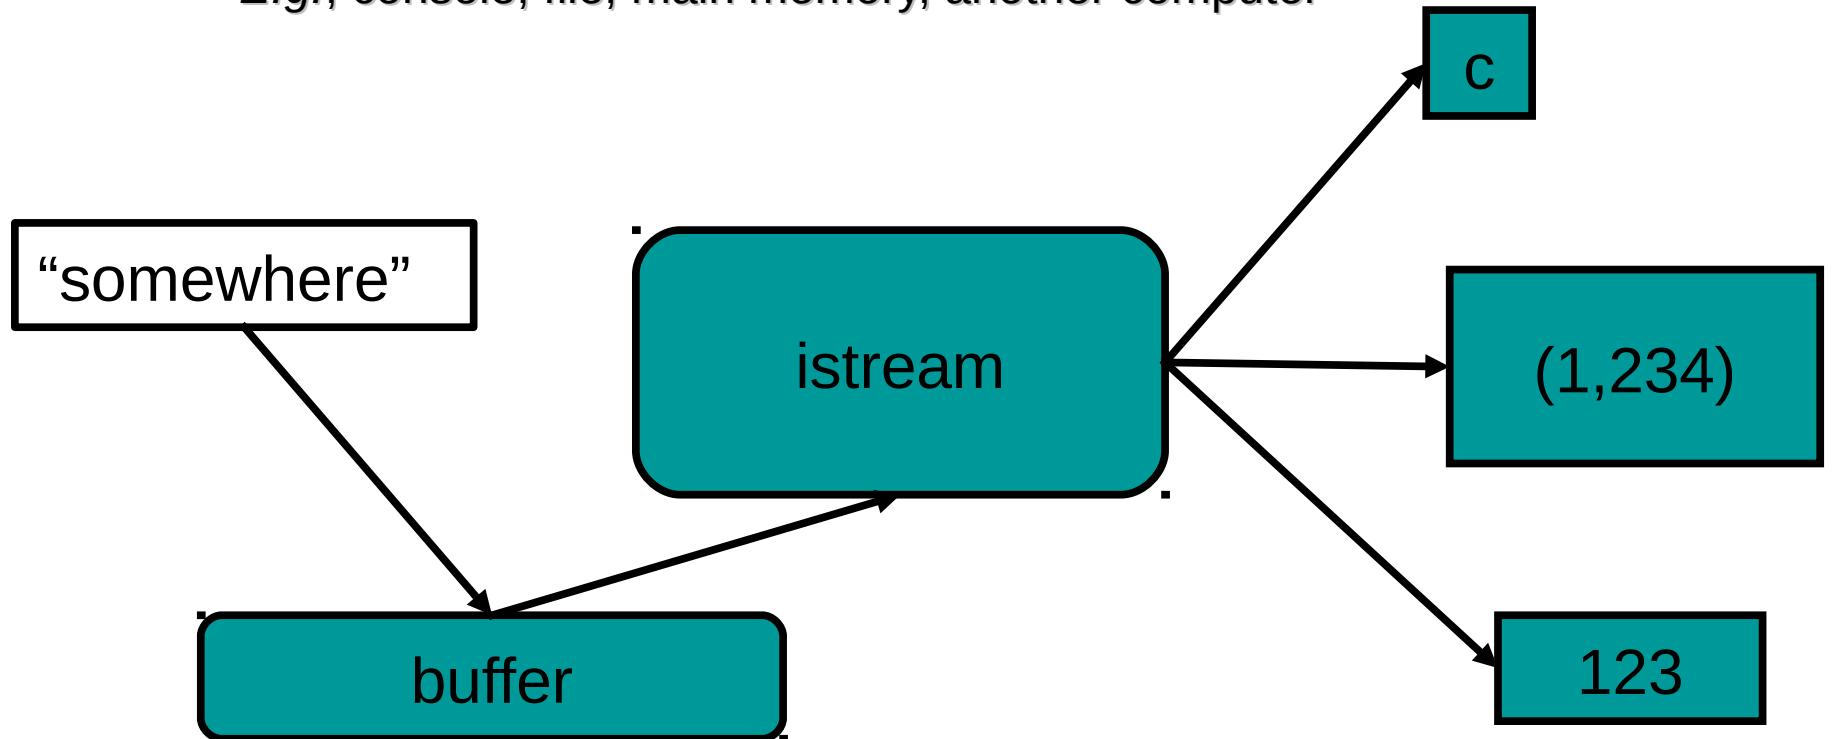
# The stream model

- An **ostream**
  - turns values of various types into character sequences
  - sends those characters somewhere
    - *E.g.*, console, file, main memory, another computer



# The stream model

- An **istream**
  - turns character sequences into values of various types
  - gets those characters from somewhere
    - *E.g.*, console, file, main memory, another computer



# Whither the Stream?

- What you do with the data in the stream determines the program flow
- Example: Handling an XML file
  - DOM: The XML stream's tags are interpreted to build a complete Document Object Model structure containing all the data in memory
  - SAX: The XML stream's tags trigger events – calls to your program's methods to handle each tag
- Both approaches are valid and will parse the XML file, but the code looks *very* different

```
<note>
<to>A. Student</to>
<from>Prof Rice</from>
<heading>Reminder</heading>
<body>Exam #3 is Apr 13.</body>
</note>
```

# The stream model

- **Reading and writing**
  - Of typed entities
    - << (output) and >> (input) plus other operations
    - Type safe
    - Formatted
  - Typically stored (entered, printed, etc.) as text
    - But not necessarily (see binary streams in chapter 11)
  - Extensible
    - You can define your own I/O operations for your own types
  - A stream can be attached to any I/O or storage device

# Files

- We turn our computers on and off
  - The contents of our main memory is transient
- We like to keep our data
  - So we keep what we want to preserve on disks and similar permanent storage
- A file is a self-contained named sequence of bytes available via the operating system
  - A file has a name
  - The file has a data format
- We can read/write a file if we know its name and format
  - Not always a given, e.g., early Microsoft Office file formats

**File** – A self-contained named sequence of bytes available via the operating system



# A File

0: 1: 2:



- At the fundamental level, a file is a sequence of bytes numbered from 0 upwards
- Other notions can be supplied by programs that interpret a “file format”
  - For example, the 6 bytes "123.45" might be interpreted as the floating-point number 123.45

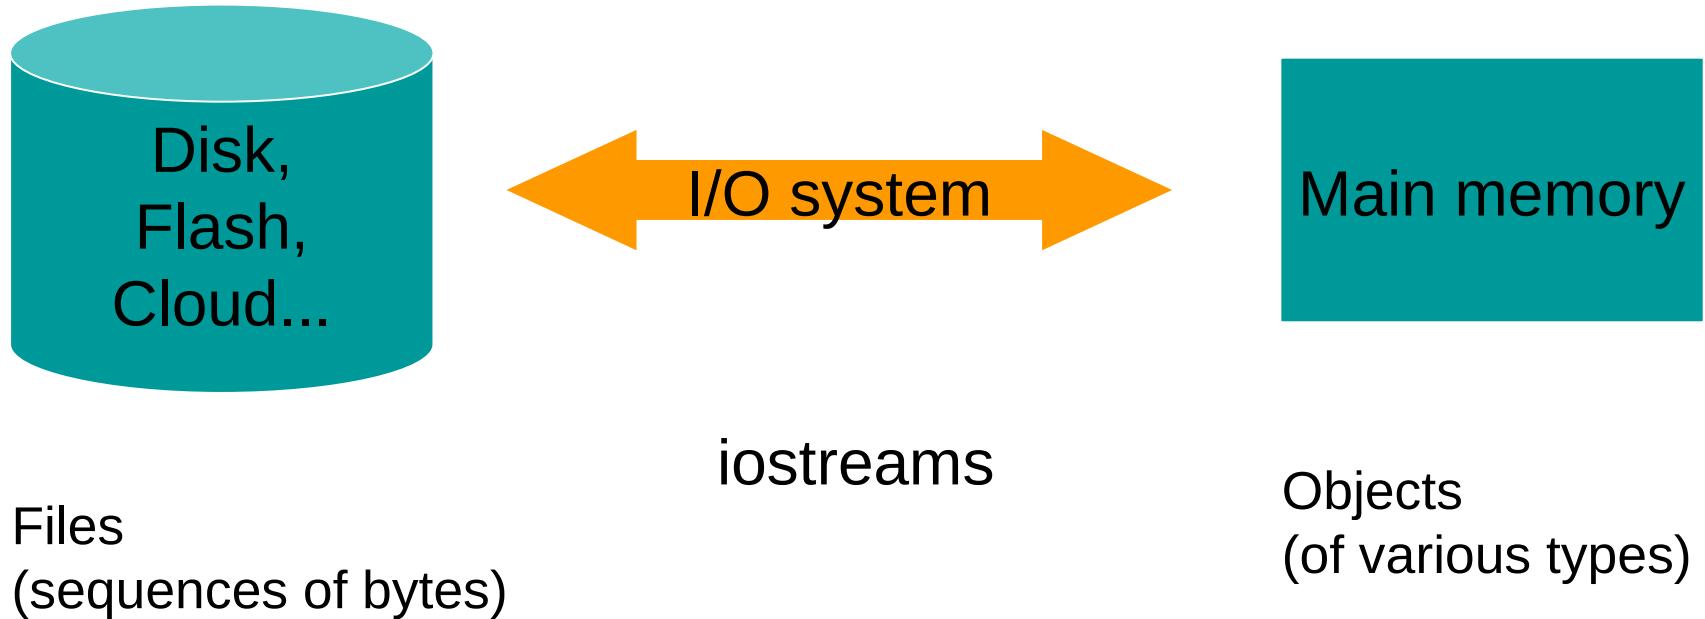
**You can examine binary files**

In Linux and Unix, “hexdump -C [file]” will reveal the contents of a binary file  
Windows requires a hex editor app or a Powershell script – GIYF



# Files

- General model



# Files

- To read a file
  - We must know its name
  - We must open it (for reading)
  - Then we can read
  - Then we must close it
    - That is typically done implicitly
- To write a file
  - We must know its name
    - Or decide on a name for a new file
  - We must open it (for writing)
    - Or create a new file of that name
  - Then we can write it
  - We must close it
    - That is typically done implicitly

# Opening a file for reading

```
#include <iostream>
#include <fstream>
#include <stdexcept>

using namespace std;

// ...
int main()
{
    cout << "Please enter input file name: ";
    string fname;
    cin >> fname;
    // The following line only works in later C++ versions, e.g., 11
    ifstream ist {fname}; // ifstream is an "input stream from a file"
                          // defining an ifstream with a name string
                          // opens the file of that name for reading
    if (!ist) throw runtime_error("can't open input file " + fname);
    // ...
```

# Opening a file for writing

```
#include <iostream>
#include <fstream>
#include <stdexcept>
using namespace std;

int main () {
    cout << "Please enter name of output file: ";
    string oname;
    cin >> oname;

    ofstream ofs {oname};
    if (!ofs) throw runtime_error("can't open output file " + oname);

    ofs << "Writing this to a file.\n";
    return 0;
}
```

# Reading from a file

- Suppose a file contains a sequence of pairs representing hours and temperature readings
  - 0 60.7
  - 1 60.6
  - 2 60.3
  - 3 59.22
- The hours are numbered 0..23
- No further format is assumed
  - Maybe we can do better than that (but not just now)
- Termination
  - Reaching the end of file terminates the read
  - Anything unexpected in the file terminates the read
    - *E.g.*, q

# Reading a File

```
#include <iostream>
#include <vector>
#include <stdexcept>
using namespace std;

class Reading { // a temperature reading
public:
    Reading(int p_hour, double p_temp)
        : hour(p_hour), temperature(p_temp) {}
    int get_hour() {return hour;}
    double get_temp() {return temperature;}
private:
    int hour;          // hour after midnight [0:23]
    double temperature;
};

int main() {
    vector<Reading> temps; // create a vector to store the readings

    int hour;
    double temperature;
    while (cin >> hour >> temperature) {                      // read
        if (hour < 0 || 23 < hour)
            throw runtime_error("hour out of range"); // check
        temps.push_back(Reading(hour, temperature)); // store
    }
}
```

# Reading a File Redux

```
#include <iostream>
#include <vector>
#include <stdexcept>
using namespace std;

class Reading { // a temperature reading
public:
    Reading(int p_hour, double p_temp)
        : hour(p_hour), temperature(p_temp) {validate();}
    int get_hour() {return hour;}
    double get_temp() {return temperature;}
    friend istream &operator>>(istream &ist, Reading &r);
    void validate() {
        if (hour < 0 || 23 < hour) throw runtime_error("hour out of range");
    }
private:
    int hour;          // hour after midnight [0:23]
    double temperature;
};

istream& operator>>(istream &ist, Reading &r) {
    ist >> r.hour >> r.temperature;
    r.validate();
    return ist;
}
int main(){
    vector<Reading> temps;           // create a vector to store the readings
    Reading reading{0,0};            // temporary to hold input
    while (cin >> reading) temps.push_back(reading); // read and store
}
```

Most logic is moved into the class  
(including data validation!)

# I/O Error Handling

- Sources of errors
  - Human mistakes
  - Files that fail to meet specifications
  - Specifications that fail to match reality
  - Programmer errors
  - Etc.
- `iostream` reduces all errors to one of four states
  - **good()** *// the operation succeeded*
  - **eof()** *// we hit the end of input ("end of file")*
  - **fail()** *// something unexpected happened (including format error)*
  - **bad()** *// something unexpected and serious happened*
- Note that `good()` and `bad()` are not opposites, despite the name
  - Name selection for these methods was especially tortured...

# Sample integer read “failure”

- Ended by “terminator character”
  - 1 2 3 4 5 \*
  - State is **fail()**
- Ended by format error
  - 1 2 3 4 5.6
  - State is **fail()**
- Ended by “end of file”
  - 1 2 3 4 5 end of file
  - 1 2 3 4 5 Control-Z (Windows)
  - 1 2 3 4 5 Control-D (Mac\* / Linux)
  - State is **eof()**
- Something really bad
  - Disk format error
  - State is **bad()**

\* A second Control-D at the start of a line may also be required

# I/O error handling

```
void fill_vector(istream& ist, vector<int>& v, char terminator)
{
    // read integers from ist into v until we reach eof() or terminator
    for (int i; ist >> i; ) // read until "some failure"
        v.push_back(i); // store in v
    if (ist.eof()) return; // fine: we found the end of file
    if (ist.bad())
        throw runtime_error("ist is bad"); // stream corrupted; let's get out of here!

    if (ist.fail()) { // clean up the mess as best we can and report the problem
        ist.clear(); // clear stream state, so that we can look for terminator
        char c;
        ist >> c; // read a character, hopefully terminator
        if (c != terminator) { // unexpected character
            ist.unget(); // put that character back
            ist.clear(ios_base::failbit); // set the state back to fail()
        }
    }
}
```

# Throwing I/O Exception on Bad()

- Request that a stream throw an exception if it goes bad

```
// How to make ist throw an exception if the stream goes bad:  
ist.exceptions(ist.exceptions() | ios_base::badbit); // add badbit to exception mask
```

- This allows us to simplify input to this

```
void fill_vector(istream& ist, vector<int>& v, char terminator)  
{    // read integers from ist into v until we reach eof() or terminator  
    for (int i; ist >> i; )  
        v.push_back(i);  
    if (ist.eof()) return;    // fine: we found the end of file  
  
    // not good() and not bad() and not eof(), ist must be fail()  
    ist.clear();    // clear stream state  
    char c;  
    ist >> c;        // read a character, hopefully terminator  
    if (c != terminator) {    // ouch: not the terminator, so we must fail  
        ist.unget(); // maybe my caller can use that character  
        ist.clear(ios_base::failbit); // set the state back to fail()  
    }  
}
```

# Reading a single value

- (At least) three kinds of problems are possible
  - the user types an out-of-range value
  - getting no value (end of file)
  - the user types something of the wrong type (here, not an integer)

```
// first simple and flawed attempt:

cout << "Please enter an integer in the range 1 to 10 (inclusive):\n";
int n = 0;
while (cin>>n) {           // read
    if (1<=n && n<=10) break; // check range
    cout << "Sorry, "
        << n
        << " is not in the [1:10] range; please try again\n";
}

// use n here
```

# Reading a single value

- What do we want to do in those three cases?
  - handle the problem in the code doing the read?
  - throw an exception to let someone else handle the problem (potentially terminating the program)?
  - ignore the problem?
- Reading a single value
  - Is something we often do many times
  - We want a solution that's very simple to use

# Handle everything: What a mess!

```
cout << "Please enter an integer in the range 1 to 10 (inclusive):\n";
int n = 0;
while (cin >> n) {
    if (cin) { // we got an integer; now check it:
        if (1<=n && n<=10) break;
        cout << "Sorry, " << n << " is not in the [1:10] range; please try again\n";
    }
    else if (cin.fail()) { // we found something that wasn't an integer
        cin.clear(); // we'd like to look at the characters
        cout << "Sorry, that was not a number; please try again\n";
        for (char ch; cin>>ch && !isdigit(ch); ) // throw away non-digits
            /* nothing */ ;
        if (!cin) error("no input"); // we didn't find a digit: give up
        cin.unget(); // put the digit back, so that we can read the number
    }
    else
        throw runtime_error("no input"); // eof or bad: give up
}
// if we get here n is in [1:10]
```

# The mess: trying to do everything at once

- Problem: We have all mixed together
  - reading values
  - prompting the user for input
  - writing error messages
  - skipping past “bad” input characters
  - testing the input against a range
- Solution: Split it up into logically separate parts

# What do we want?

- What logical parts do we want?
  - **int get\_int(int low, int high);** *// read an int in [low..high] from cin*
  - **int get\_int();** *// read an int from cin*  
*// so that we can check the range int*
  - **void skip\_to\_int();** *// we found some “garbage” character*  
*// so skip until we find an int*
- Separate functions that do the logically separate actions

# Skip “garbage”

```
void skip_to_int()
{
    if (cin.fail()) {      // we found something that wasn't an integer
        cin.clear(); // we'd like to look at the characters
        for(char ch; cin>>ch; ) { // throw away non-digits
            if (isdigit(ch) || ch=='-') {
                cin.unget(); // put the digit back,
                    // so that we can read the number
                return;
            }
        }
    }
    error("no input");    // eof or bad: give up
}
```

# Get (any) integer

```
int get_int()
{
    int n = 0;
    while (true) {
        if (cin >> n) return n;
        cout << "Sorry, that was not a number; please try again\n";
        skip_to_int();
    }
}
```

# Get integer in range

```
int get_int(int low, int high)
{
    cout << "Please enter an integer in the range "
        << low << " to " << high << " (inclusive):\n";
    while (true) {
        int n = get_int();
        if (low<=n && n<=high) return n;
        cout << "Sorry, "
            << n << " is not in the [" << low << ':' << high
            << "] range; please try again\n";
    }
}
```

# Use

- Problem:
  - The “dialog” is built into the read operations

```
int n = get_int(1,10);
cout << "n: " << n << endl;

int m = get_int(2,300);
cout << "m: " << m << endl;
```

# What do we *really* want?

- That's often the really important question
- Ask it repeatedly during software development
- As you learn more about a problem and its solution, your answers improve

```
// parameterize by integer range and "dialog"

int strength = get_int(1, 10,
                      "enter strength",
                      "Not in range, try again");
cout << "strength: " << strength << endl;

int altitude = get_int(0, 50000,
                      "please enter altitude in feet",
                      "Not in range, please try again");
cout << "altitude: " << altitude << "ft. above sea level\n";
```

[Note: Separating the control from the view is even better]

# Parameterize

- Incomplete parameterization: `get_int()` still “blabbers”
  - “utility functions” should not produce their own error messages
  - Serious library functions do not produce error messages at all
    - They throw exceptions (possibly containing an error message)

```
int get_int(int low, int high, const string& greeting, const string& sorry)
{
    cout << greeting << ":" [ " << low << ':' << high << "]\n";
    while (true) {
        int n = get_int();
        if (low<=n && n<=high) return n;
        cout << sorry << ":" [ " << low << ':' << high << "]\n";
    }
}
```

# User-defined output: operator<<()

- Usually trivial
- We often use several different ways of outputting a value
  - Tastes for output layout and detail vary

```
ostream& operator<<(ostream& os, const Date& d)
{
    return os << '(' << d.year()
        << ',' << d.month()
        << ',' << d.day() << ')';
}
```

# A Note on \*this or, Making the Sausage

```
void do_some_printing(Date d1, Date d2)
{
    cout << d1;                  // means operator<<(cout,d1) ;

    cout << d1 << d2;
        // means (cout << d1) << d2;
        // means (operator<<(cout,d1)) << d2;
        // means operator<<((operator<<(cout,d1)), d2) ;
}
```

# A Note on \*this or, Making the Sausage

- When calling a method, C++ creates a pointer to the method's object called `*this`
    - So, in `cout << d1;`, `*this` points to `cout`
    - C++ sees the above code as  
`operator<<(cout, d1);`
    - The method returns `*this`, which points to `cout`
  - Thus when chaining stream operators,  
e.g., `cout << d1 << d2;`, we get:
    - `(cout << d1) << d2;`
    - `(operator<<(cout, d1)) << d2;` which is `(cout) << d2;`
- or, fully nesting the operations
- `operator<<(operator<<(cout, d1)), d2);`

# User-defined input: operator>>()

```
istream& operator>>(istream& is, Date& dd)
    // Read date in format: ( year , month , day )
{
    int y, m;
    char ch1, ch2, ch3, ch4;
    is >> ch1 >> y >> ch2 >> m >> ch3 >> d >> ch4;
    if (!is) return is; // we didn't get our values, so just leave
    if (ch1 != '(' || ch2 != ',' || ch3 != ',' || ch4 != ')') { // oops: format error
        is.clear(ios_base::failbit); // something wrong: set state to fail()
        return is; // and leave
    }
    dd = Date{y, Month(m), d}; // update dd
    return is; // and leave with is in the good() state
}
```

# Quick Review

- A named sequence of bytes available via the operating system is called a \_\_\_\_\_.
- To parse a file, we must know its \_\_\_\_\_ and \_\_\_\_\_ \_\_\_\_\_.
- Throwing \_\_\_\_\_ reports a generic error that occurred during the time the program was running. The parameter is a \_\_\_\_\_ that describes the error.
- When reading from a stream, what do these statuses mean?
  - good()
  - eof ()
  - bad()
  - fail()
- The end of file keystroke is ^\_\_ in Windows and ^\_\_ in Mac / Linux.
- The \_\_\_\_\_ method of the stream allows us to specify a mask enabling the throwing of exceptions for desired stream events.

# For Next Class

- (Optional) Read Chapter 10 in Stroustrup
  - Do the Drills!
- Skim Chapter 11 for next lecture
- **Sprint #2 now in progress: It's GUI Week!**
  - **Individuals:** Create your main window, and dialogs for instancing a serving (and a serving class)
  - **Duos:** Also create dialogs for instancing servers and customers, and to create orders (and an order class)
  - **Trios:** Also manage the state of the orders, and create the Emporium itself (and an emporium class)
  - **Quattros:** Also display orders for servers and customers, and load and save data to files