

**CSE 1325: Object-Oriented Programming**

**Lecture 19**

# **Embedded Programming, and UML State Diagrams with a C++ Implementation**

**Mr. George F. Rice**

**[george.rice@uta.edu](mailto:george.rice@uta.edu)**

**Based on material by Bjarne Stroustrup**

**[www.stroustrup.com/Programming](http://www.stroustrup.com/Programming)**

**ERB 402**

**Office Hours:**

**Tuesday Thursday 11 - 12**

**Or by appointment**

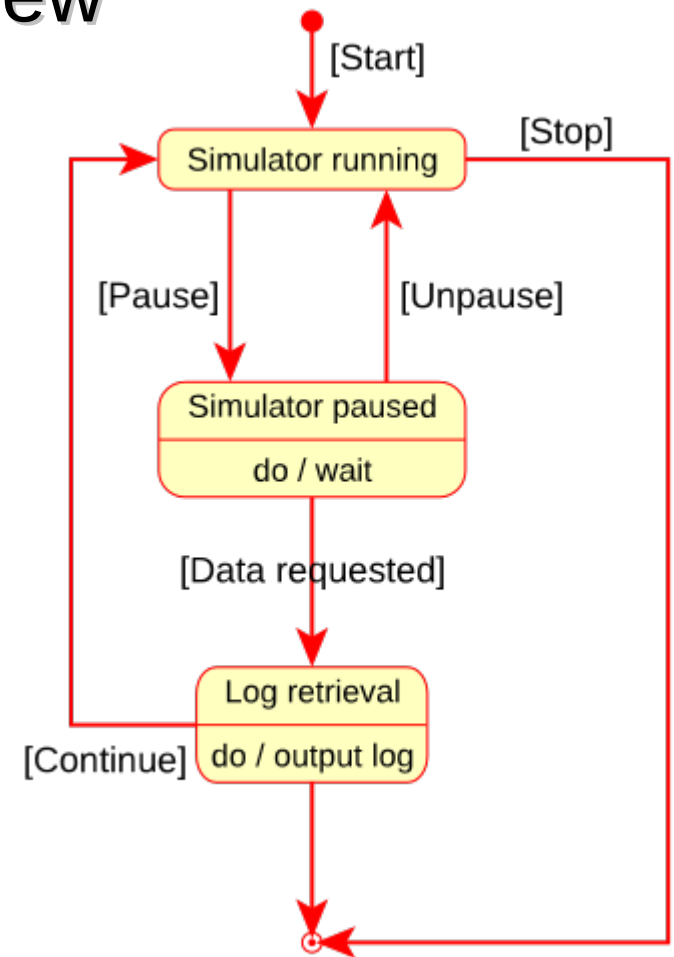


# Quick Review

- **True** or False: Most (but not all) stream operators are “sticky”. If False, which is it? If True, give an example of each. **Hex is sticky, setw is not sticky.**
- True or **False**: Stream operators are constants and thus cannot accept parameters.
- To output hexadecimal numbers via cout, include the **hex** operator in the stream. To precede hexadecimal numbers with “0x”, include **showbase** in the stream.
- What happens if a value exceeds the specified stream output width? **Correct data takes precedence over formats such as output width**
- **True** or False: Binary file operations are inherently less portable than text file operations.
- Stream operations can target string variables by using the **stringstream** class.
- The **Strategy** pattern enables an algorithm’s behavior to be modified at runtime.
- The UML Activity diagram displays a sequence of activities at the **algorithm** level, particularly useful for documenting Use Cases.
- Which of the following is supported by the UML Activity Diagram?  
(a) Friends **(b) Interrupts** **(c) Swimlanes** **(d) Hierarchical Diagrams** (e) Inheritance

# Overview: UML State Diagrams

- Embedded Programming Overview
- UML State Diagrams
  - Elements
  - Mealy and Moore
  - Hierarchical State Machines
- State Design Pattern
- C++ Realization
  - Adding Event Handling
  - Dealing with Forward Refs
  - Testing






# Embedded Programming

- Embedded programming comes in 2 flavors
  - Hard real-time – A correct answer after the deadline is the wrong answer
  - Soft real-time – A correct answer should be delivered before the deadline most of the time
- The software runs on special hardware
  - Hardware issues must be explicitly handled
  - Portability to next-gen hardware must be pre-planned
  - This is sometimes the bulk of the code



# Embedded Systems



- Computers used as part of a larger system
  - That usually don't look like a computer
  - That usually control physical devices
- Often reliability is critical
  - “Critical” as in “if the system fails someone might die”
- Often resources (memory, processor capacity) are limited
- Often real-time response is important – or essential



# Examples of Embedded Systems

- Assembly line quality monitors
- Bar code readers
- Bread machines
- Cameras
- Car assembly robots
- Cell phones
- Centrifuge controllers
- CD players
- Disk drive controllers
- “Smart card” processors
- Fuel injector controls
- Medical equipment monitors
- PDAs
- Printer controllers
- Sound systems
- Rice cookers
- Telephone switches
- Water pump controllers
- Welding machines
- Windmills
- Wrist watches





# Who Works on Embedded Systems?

- Computer Scientists
  - Unless you only do web and client-server
- Computer Engineers
  - Hardware and software are complementary aspects of the overall system
  - The hardware must support the software well
- Electrical Engineers
  - You must be able to talk to the computer guys
  - Your concerns are also complementary – RF, power, feedback control systems, etc.






# Hard Real-Time Uses a (Large) Subset of C++

- Execution time must be highly predictable
  - Memory allocation (new) times are highly variable
    - Depends on how fragmented heap is currently
  - Exception times are difficult to pin down
    - Execution path moves non-linearly up the call stack
  - System libraries must be included with caution
    - A single include can drag in a bloating of code
  - State machines often model real-world problems well – if the implementation is efficient and predictable
- Measure relentlessly and know your environment!



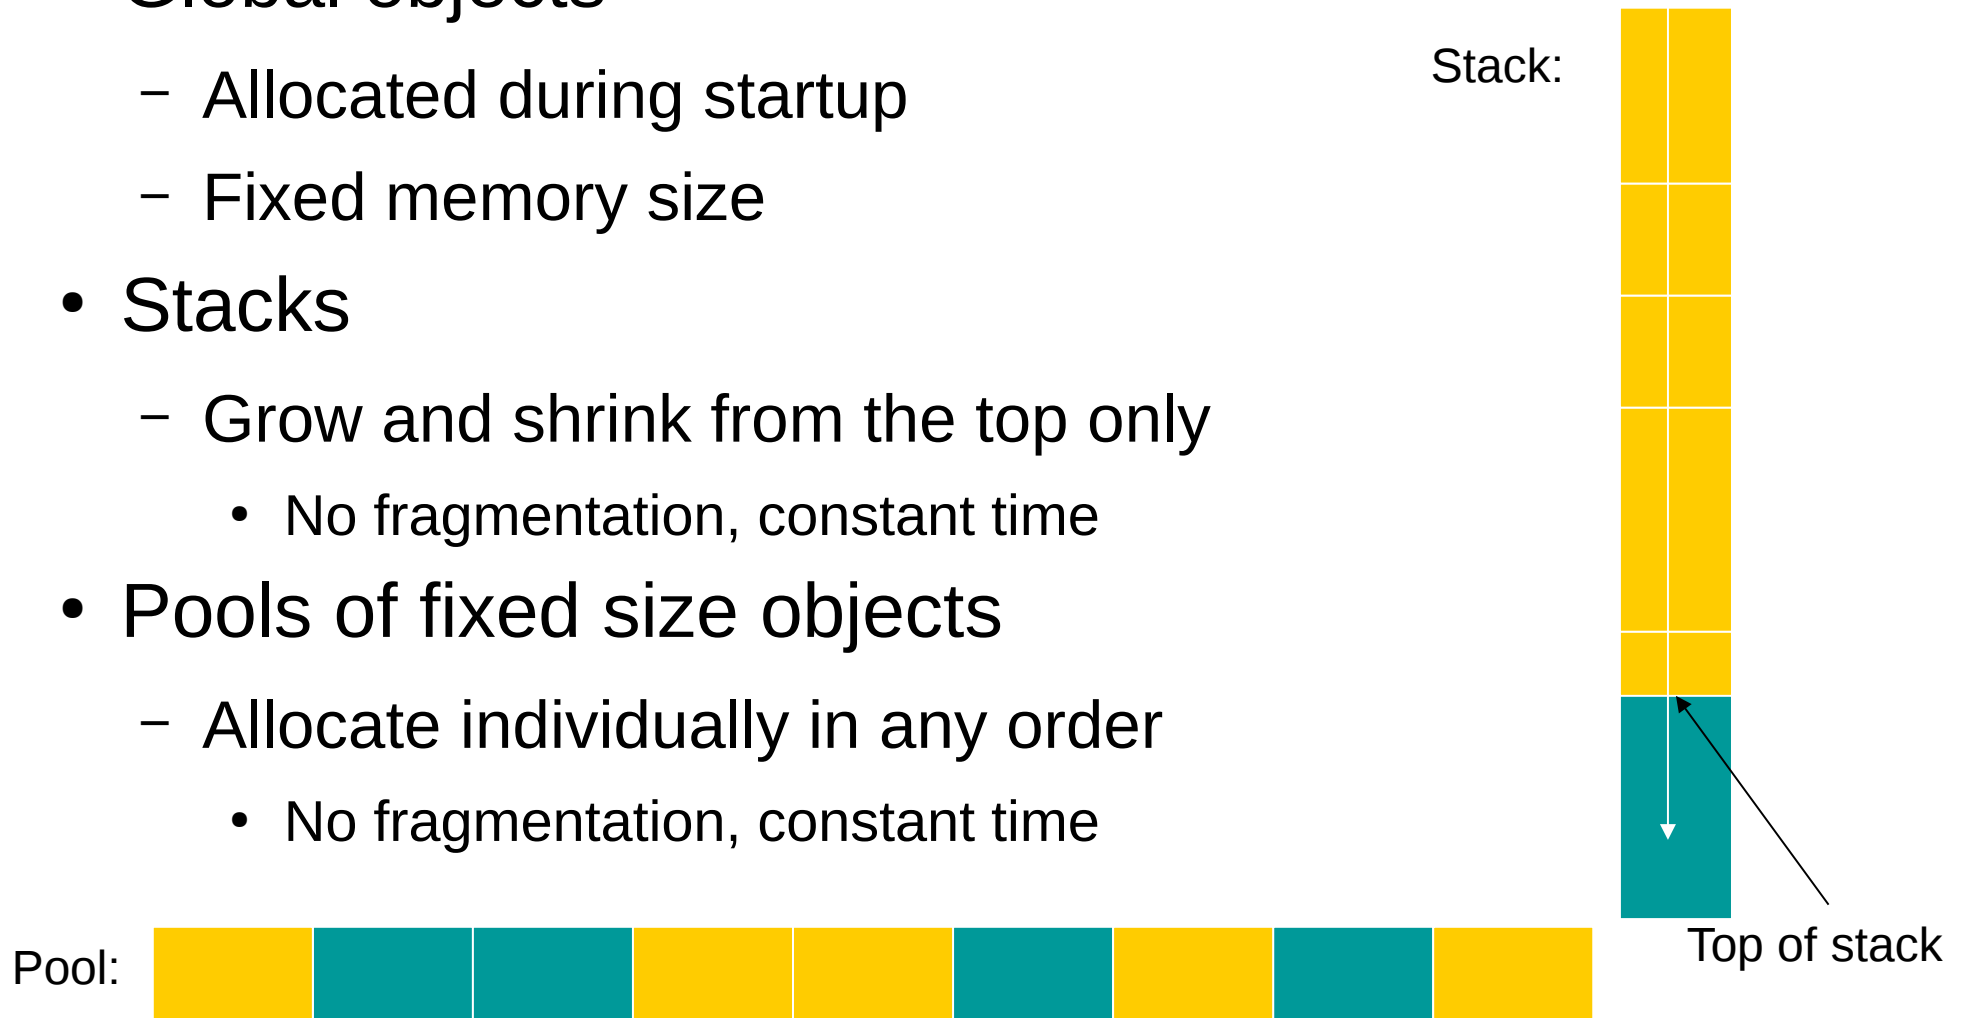
# The Trouble with New and Delete

- C++ code refers directly to memory
  - Once allocated, an object cannot be moved (or can it?)
- Allocation delays
  - The effort needed to find a new free chunk of memory of a given size depends on what has already been allocated
- Fragmentation
  - If you have a “hole” (free space) of size  $N$  and you allocate an object of size  $M$  where  $M < N$  in it, you now have a fragment of size  $N - M$  to deal with
  - After a while, such fragments constitute much of the memory



# The Solution is Preallocated Structures

- Global objects
  - Allocated during startup
  - Fixed memory size
- Stacks
  - Grow and shrink from the top only
    - No fragmentation, constant time
- Pools of fixed size objects
  - Allocate individually in any order
    - No fragmentation, constant time





# Pool Example

```
// Note: element type is known at compile time
// Allocation times are completely predictable (and short)
// The user has to pre-calculate the maximum number of elements needed

template<class T, int N> class Pool {
public:
    Pool();           // make pool of N Ts - construct pools only during startup
    T* get();         // get a T from the pool; return 0 if no free Ts
    void free(T* p);  // return a T given out by get() to the pool
private:
    // keep track of T[N] array (e.g., a list of free objects)
};

// Example use
Pool<Small_buffer, 10> sb_pool;
Pool<Status_indicator, 200> indicator_pool;
```

# Stack Example

```
// Note: allocation times completely predictable (and short)
// The user has to pre-calculate the maximum number of elements needed

template<int N> class Stack {
public:
    Stack(); // make an N byte stack - construct stacks only during startup
    void* get(int N); // allocate n bytes from the stack; return 0 if no free space
    void free(void* p); // return last block returned by get() to the stack
private:
    // keep track of an array of N bytes (e.g. a top of stack pointer)
};

// Examples

Stack<50*1024> my_free_store; // 50K worth of storage to be used as a stack

void* pv1 = my_free_store.get(256 * sizeof(int)); // allocate array of ints
int* pi = static_cast<int*>(pv1); // convert memory to objects

void* pv2 = my_free_store.get(50);
Pump_driver* pdriver = static_cast<Pump_driver*>(pv2);
```





# Templates Rock for Real-Time!

- Vector class (for example) is a template
  - But not one that is suitable for real-time!
- Operations are inline
  - No run-time overhead
- No memory consumed for unused operations
  - Memory is often in short supply
- Reminiscent of the preprocessor
  - Now with class(es)!

# Bit Representation in C++

- unsigned char uc; // 8 bits
- unsigned short us; // typically 16 bits
- unsigned int ui; // typically 16 bits or 32 bits
  - // (check before using)
  - // many embedded systems have 16-bit ints
- unsigned long int ul; // typically 32 bits or 64 bits
- std::vector<bool> vb(93); // 93 bits – C++ 14 only
  - true/false auto-converts to/from 1/0
  - Use only if you really need more than 32 bits
- std::bitset bs(314); // 314 bits
  - Use if you really need more than 32 bits
  - Typically efficient for multiples of sizeof(int)



# Bit Manipulation

a: 

1	0	1	0	1	0	1	0
---	---	---	---	---	---	---	---

 0xaa

b: 

0	0	0	0	1	1	1	1
---	---	---	---	---	---	---	---

 0x0f

a&b: 

0	0	0	0	1	0	1	0
---	---	---	---	---	---	---	---

 0x0a

a|b: 

1	0	1	0	1	1	1	1
---	---	---	---	---	---	---	---

 0xaf

a^b: 

1	0	1	0	0	1	0	1
---	---	---	---	---	---	---	---

 0xa5

a<<1: 

0	1	0	1	0	1	0	0
---	---	---	---	---	---	---	---

 0x54

b>>2: 

0	0	0	0	0	0	1	1
---	---	---	---	---	---	---	---

 0x03

~b: 

1	1	1	1	0	0	0	0
---	---	---	---	---	---	---	---

 0xf0

& and

(often checks a bit)

| inclusive or

(often sets a bit)

^ exclusive or

(often flips a bit;  
also for graphics, crypto)

<< left shift

>> right shift

~ one's complement



# Know Your Environment

- The disassembler is the embedded programmer's best friend
  - What code did that C++ feature generate?
  - *Why* did it generate that code?
  - What alternate features generate “better” code?
- The system library source code is for bedtime reading
  - Which classes use unacceptable features?
  - Which interdependencies exist and how do they affect your memory and latency?
- The Real-Time Operating System (RTOS) is a key foundation
  - Especially the scheduler – know it well!
- The “non-hosted environment” (remote debugger) drives your productivity during integration to a surprising degree
  - Many “bugs” are *hardware* bugs, unlike in IT!



# Failing Hardware


- Hardware failures are much more common in embedded than in IT
  - Power surges and sags
  - Connectors vibrate loose
  - Physical damage
  - Environmental extremes
- The system may be remote when it fails
  - Like, on Mars... or in the Mariana Trench



- Strategies
  - Self-check
    - Heartbeat / Reset
    - Shuttle - 4+1
  - Redundancy
    - Fly by wire
    - Shuttle - 4+1
  - Degraded Modes
    - F-16 – 8 CPUs, 256 modes
  - Debug-only bus
    - I<sup>2</sup>C (Inter-IC), step pins

# Modeling the World in States

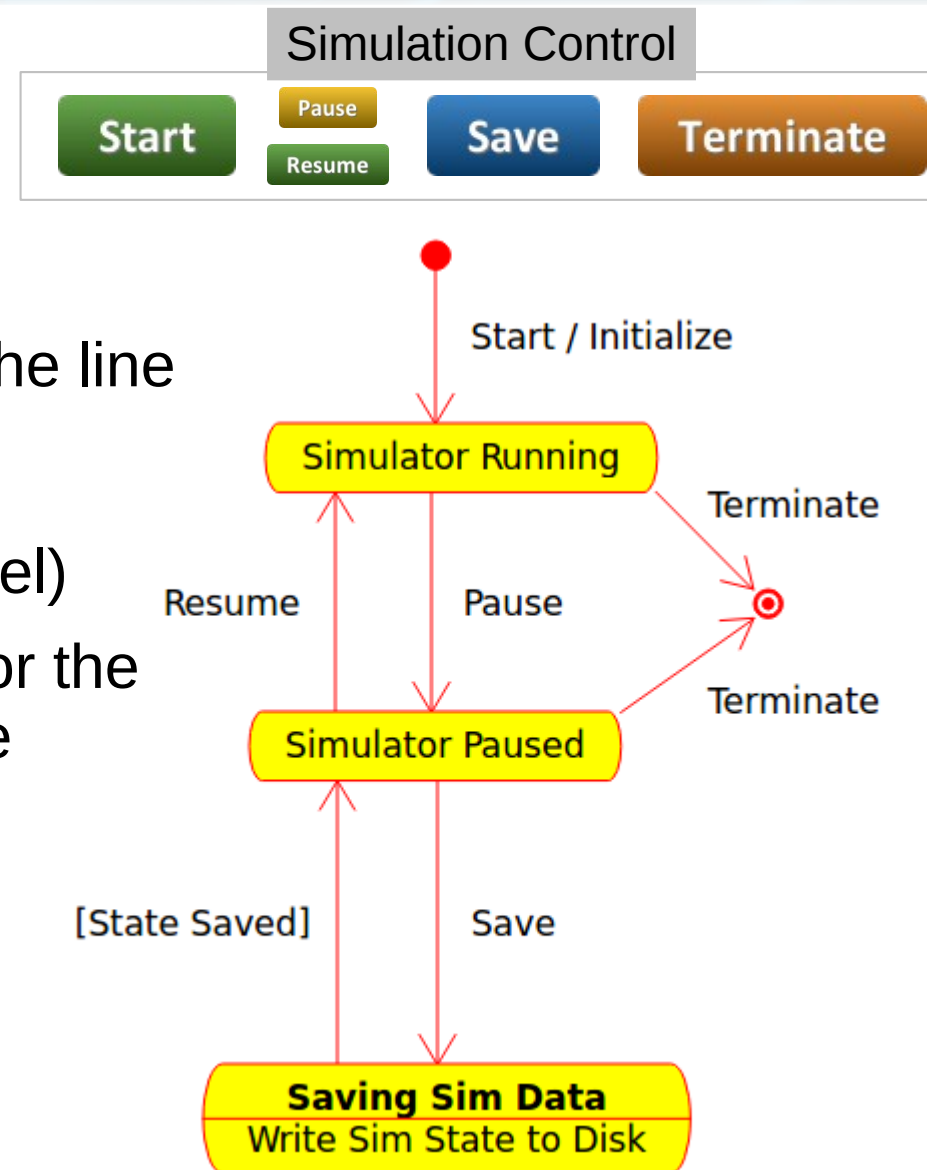
- Many embedded (and IT) systems exhibit state-like behavior
- A **State** is the cumulative value of all relevant stored information to which a system or subsystem has access.
  - The output of a stateful system is completely determined by its current state and its current inputs
  - A **State Diagram** documents the states, permissible transitions, and activities of a system






# Basic UML State Diagram Elements

- Initial state (filled circle)
- State (rounded rectangle)
  - Name of state above the line
  - Optional activity / output below the line
- Transition (arrow line)
  - Event causing the transition (label)
  - Guard (bool) that must be true for the transition to occur (inside square brackets)
  - Activity / output during transition (after the / )
- Final state (target)




# (Non-UML) Mealy vs Moore Diagrams



A Mealy state machine defines outputs based on current state and current inputs

- Outputs are defined on *transitions*
- Outputs follow inputs immediately
- Often require fewer states

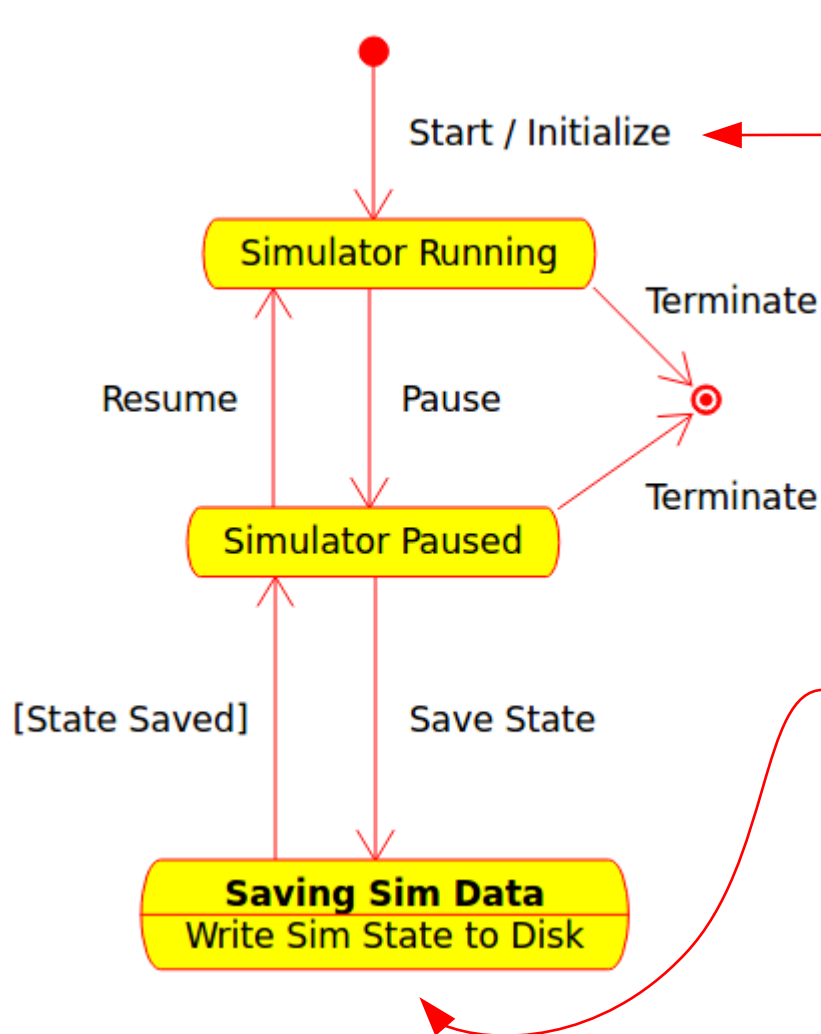


A Moore state machine defines outputs based solely on the current state

- Outputs are defined in *states*
- Outputs change only on clock edges
- Safer for hardware realization due to greater immunity from race conditions and line noise



# UML State Diagrams are Simultaneously Mealy and Moore



A Mealy state machine defines outputs based on current state and current inputs

- Outputs are defined on *transitions*
- Outputs follow inputs immediately
- Often require fewer states

A Moore state machine defines outputs based solely on the current state

- Outputs are defined in *states*
- Outputs change only on clock edges
- Safer for hardware realization due to greater immunity from race conditions and line noise



# Isn't This Just an Activity Diagram?

- Not exactly, but they are closely related
  - **An activity diagram is a special case of a state diagram**, in which the states consist only of activities and transitions are (almost) always triggered by completion of those activities rather than external events
  - Conceptually, activity diagrams may imply a data flow graph, while a state diagram may imply a structured, centralized control scheme
- Action and Activity states are conceptually different
  - An action state in a state diagram cannot be further decomposed, and happens instantly (although an associated activity may consume time)
  - An activity state in an activity diagram can be further decomposed, and the state has a definite timespan before the next transition
- Either may be used to document a use case or to model the behavior of a class or subsystem
  - As may a Sequence Diagram



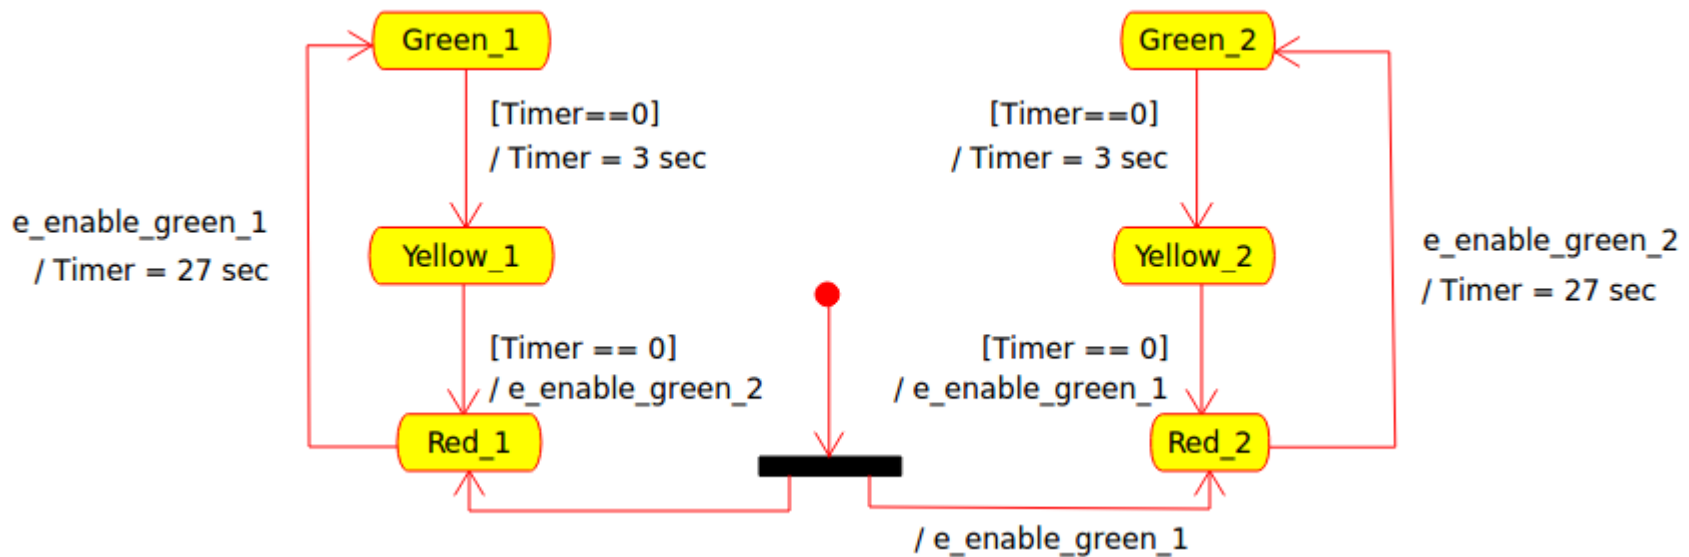


# History of UML State Diagrams

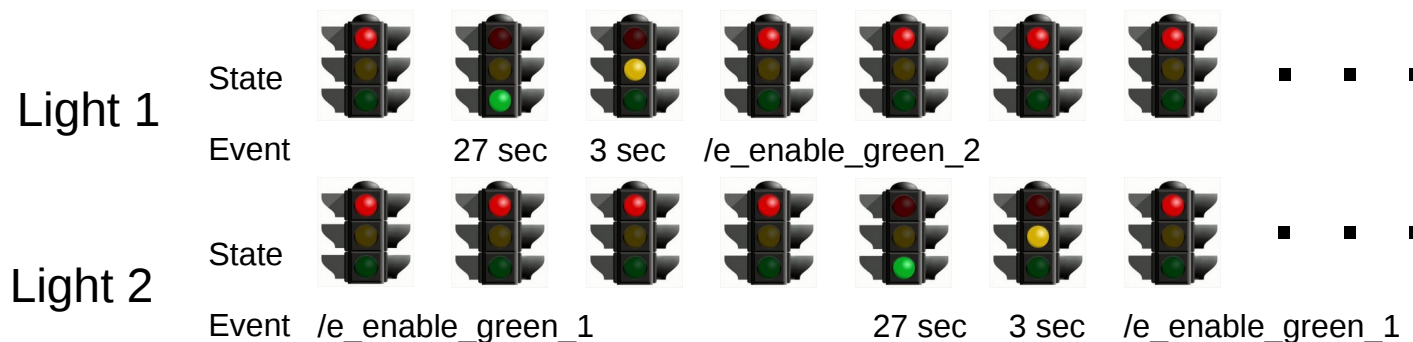
- Basic state diagrams have been around since the dawn of computers
  - States are fundamental to all clock-driven circuits
  - Hardware implementations have special challenges
- Dr. David Harel formalized and greatly extended hierarchical statechart semantics<sup>1</sup> in 1986, and is the “Father of Modern Statecharts”
- UML adopted an object-oriented variant of Harel Statecharts in 1997, and significantly extended the semantic model again

<sup>1</sup> <http://www.wisdom.weizmann.ac.il/~harel/SCANNED.PAPERS/Statecharts.pdf>

# Modeling a Traffic Light



- Each direction is controlled by a separate state machine.
- The state of each machine selects the lights in that direction.
- The state machines share a timer and communicate via events.







# States

- Each state models a distinct aggregation of *relevant* memory and hardware in the system
  - Most of memory isn't relevant to the state machine
- UML states are *extended* states, allowed to contain complex data structures to complement the state itself
  - These should be used minimally to manage complexity of the state machine
  - Extended state variables raise the need for guards

# Guards


- A **guard** is a Boolean expression that enables a state transition when true and disables it when false, e.g., [power == on]
  - The Boolean expression can reference any available object, but typically depends on related state machines, event parameters, or simple external signals
  - Guards should be minimized, as overuse can make a state machine and its associated code “brittle” and difficult to debug





# Events

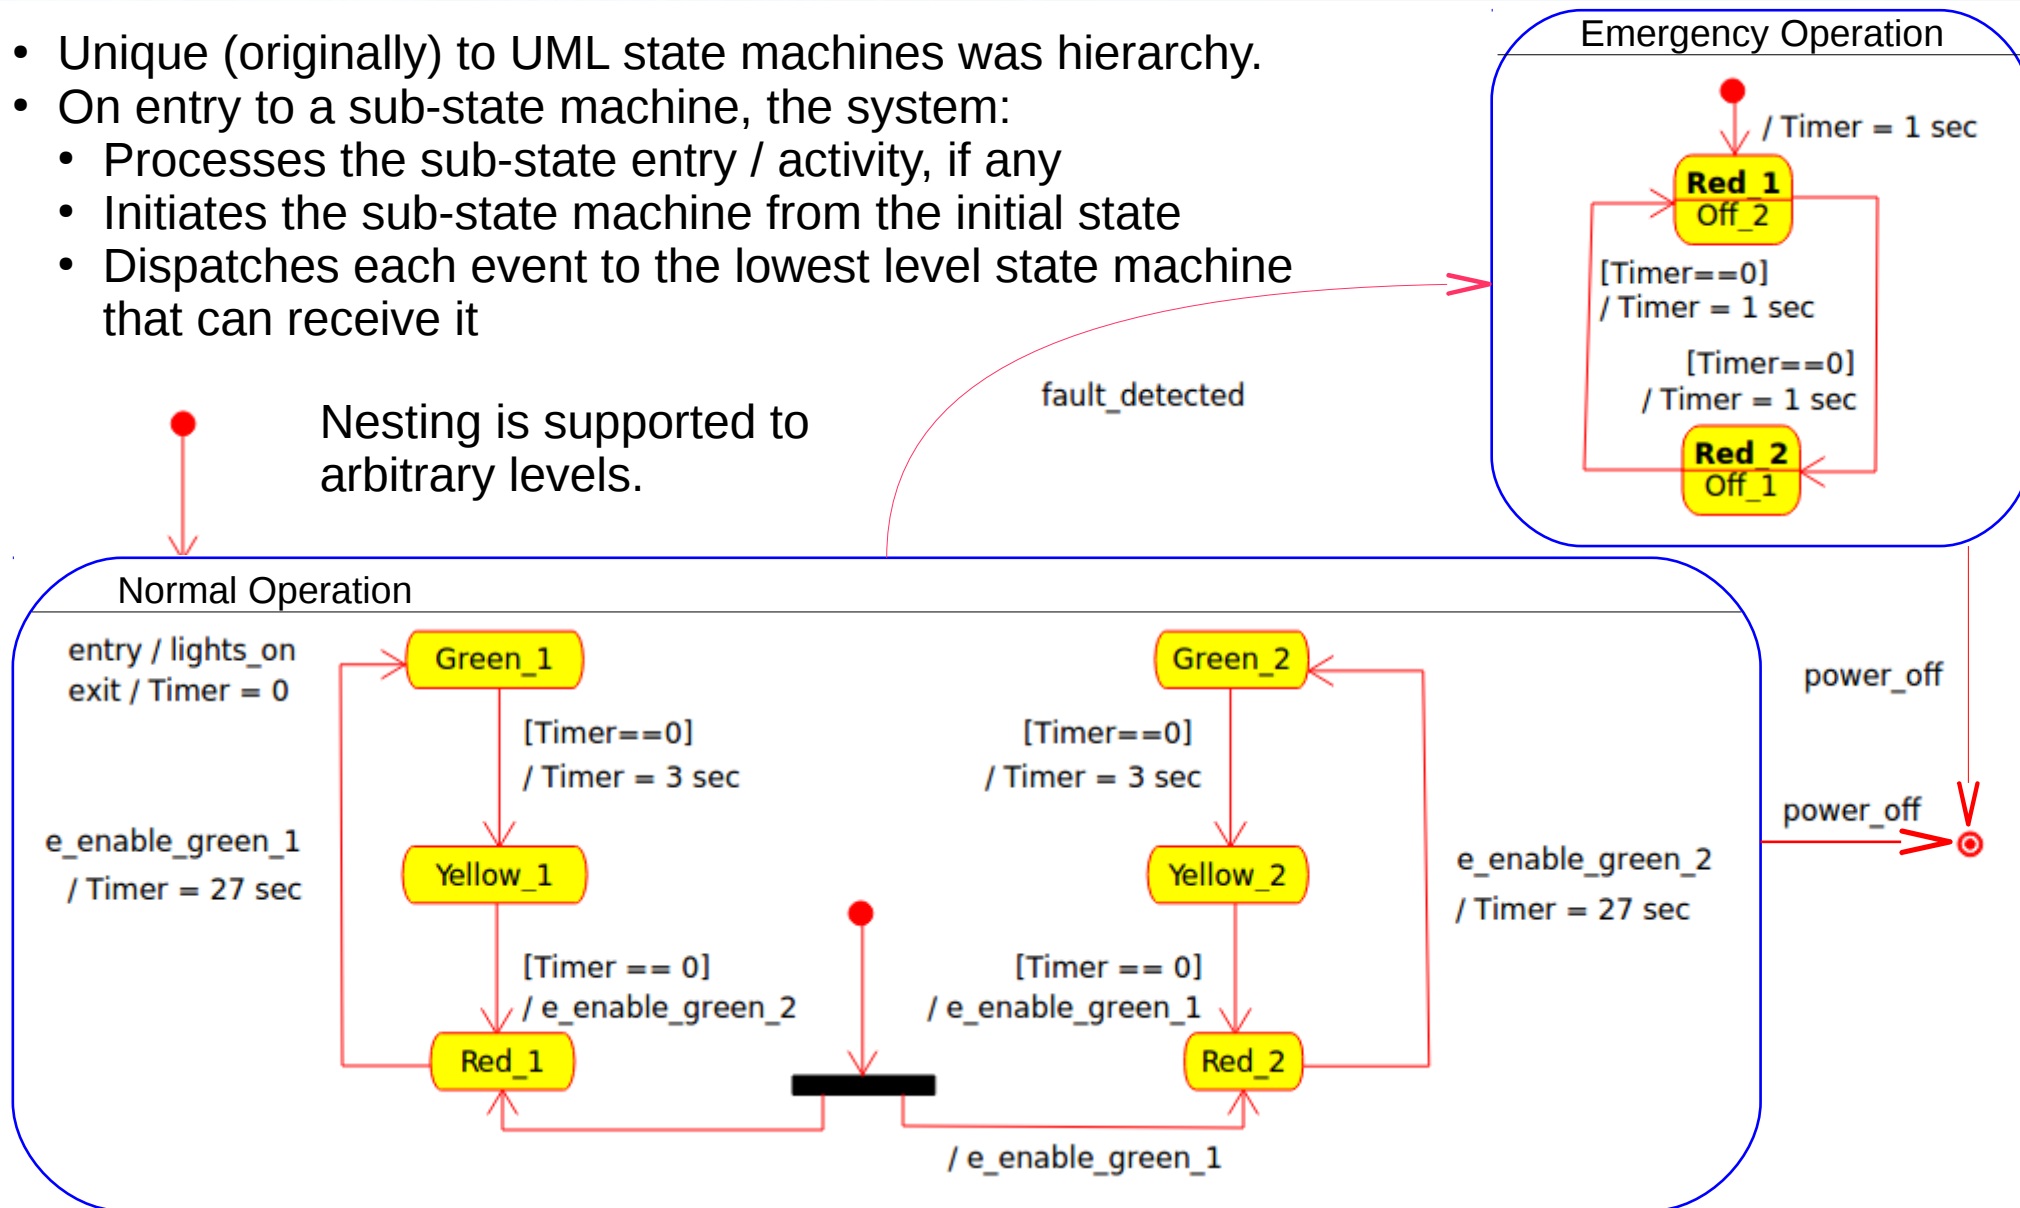
- An **event** is a type of occurrence that potentially affects the state of the system
  - The event may be generated e.g., by user action, timer expiration, or changes external to the system
  - The event may have one or more parameters
  - Events follow a 3-state lifecycle
    - Received, where it waits on an event queue until a machine reaches a state able to dispatch that event
    - Dispatched, where it affects the state machine
    - Consumed, where it is no longer available for use
  - Only one event may be dispatched at a time
  - An event which no machine can handle is quietly discarded



# Hierarchical State Machines (HSM)

- Unique (originally) to UML state machines was hierarchy.
- On entry to a sub-state machine, the system:
  - Processes the sub-state entry / activity, if any
  - Initiates the sub-state machine from the initial state
  - Dispatches each event to the lowest level state machine that can receive it

Nesting is supported to arbitrary levels.






## Structural

# State Design Pattern

- The State Design pattern supports full encapsulation of unlimited states within a scalable context
  - This is a variation on the Strategy Pattern, optimized for state-dependent behaviors
  - The pattern allows the context (state machine) behavior to depend on the currently active State instance

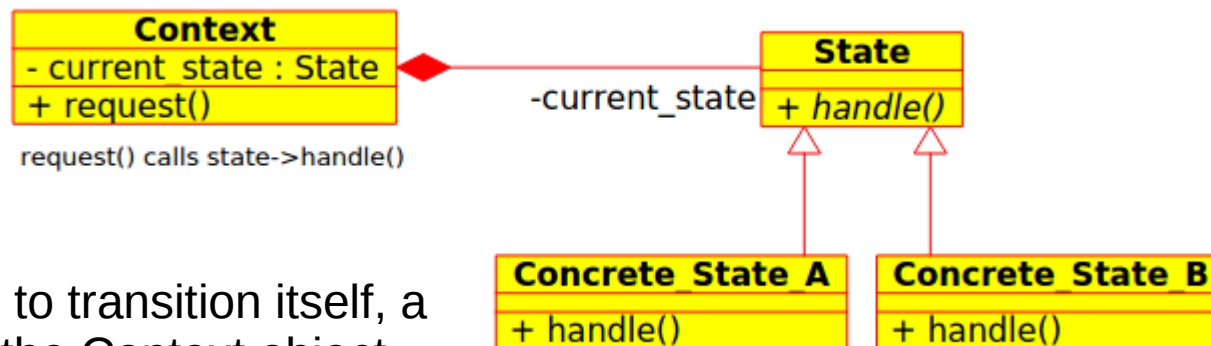




# State Design Pattern

Context is the state machine.  
It's current state is an instance  
of a concrete state class.

State is the virtual base class that declares  
(and in some cases defines) common state  
operations (represented here by *handle()*).



For the state to transition itself, a  
reference to the Context object  
must be passed to `handle()`.

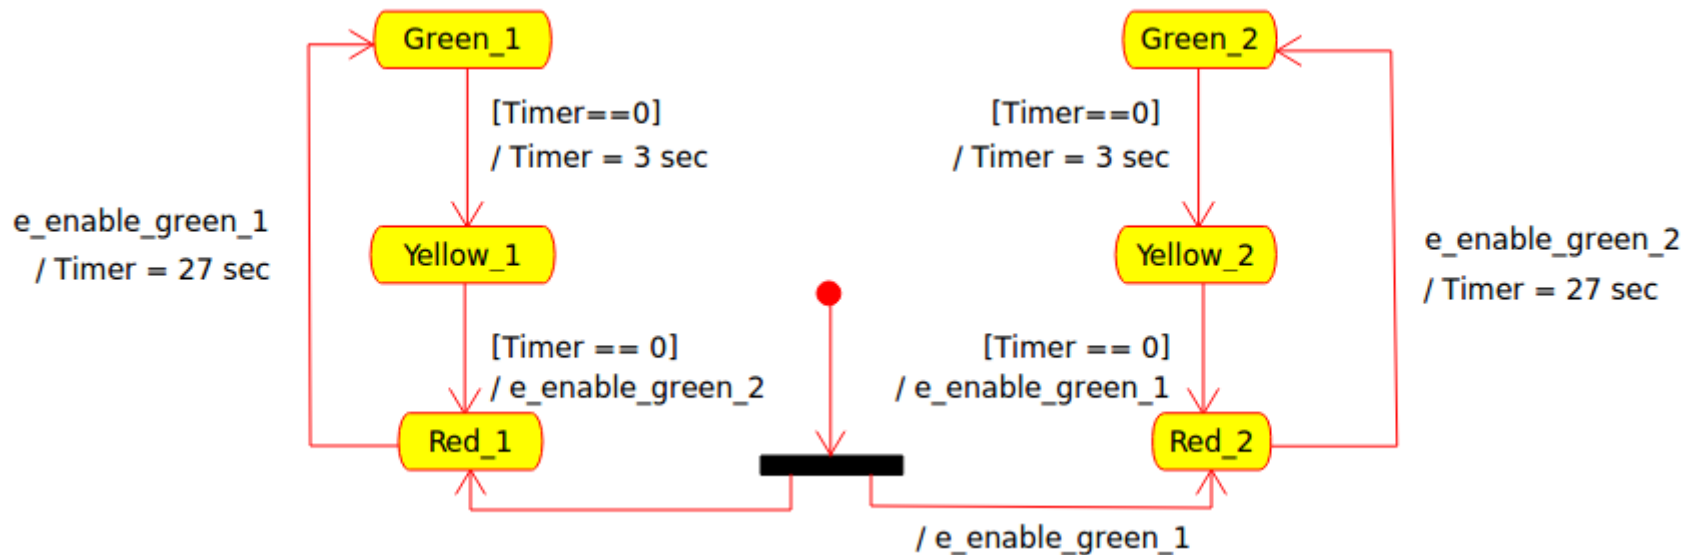
**Note that event handling is not  
addressed by this pattern**, and  
must be added to realize many  
UML state machine designs.

Classes derived from State represent the states  
defined within the state machine model. The `handle()`  
methods



# Implementing our Traffic Light

- We'll use the State Design Pattern, augmented by a basic event handler, to automate the traffic state diagram below



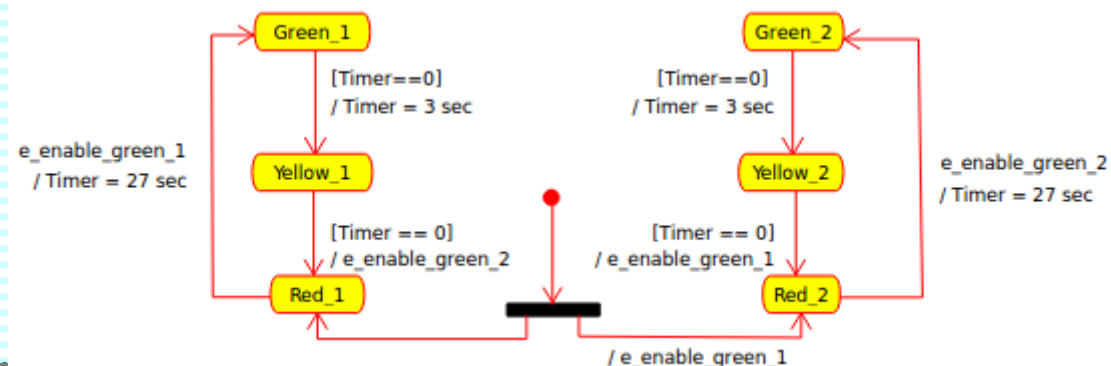

# First, We Need Light Colors

```
#include <stdexcept>
#include <iostream>

using namespace std;

//
// Traffic light colors
//

enum class Traffic_light_color {GREEN, YELLOW, RED};
string ctos(Traffic_light_color color) {
    if (color == Traffic_light_color::GREEN) return "green";
    if (color == Traffic_light_color::YELLOW) return "yellow";
    if (color == Traffic_light_color::RED) return "red";
    throw runtime_error("ctos: invalid color");
}
```



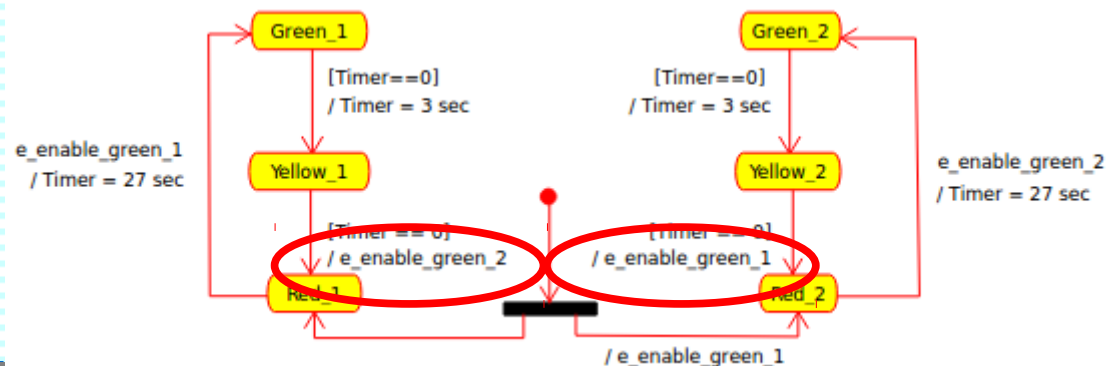



# Next, a Simple Event Handler

```
//  
// Events  
//  
  
class Event {  
public:  
    void generate() {++_pending;}  
    bool consume() {  
        if (_pending > 0) {--_pending; return true;}  
        return false;  
    }  
private:  
    int _pending = 0;  
};
```

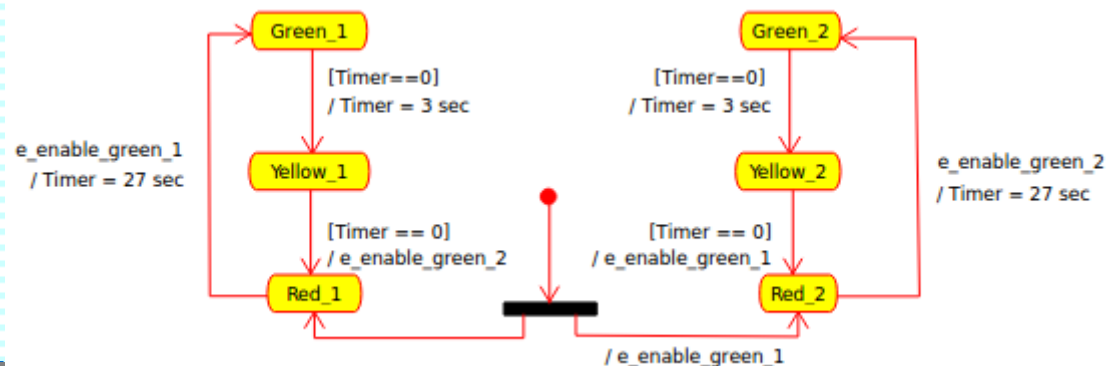

Event e\_enable\_green\_1;

Event e\_enable\_green\_2;



# Our Virtual State Class

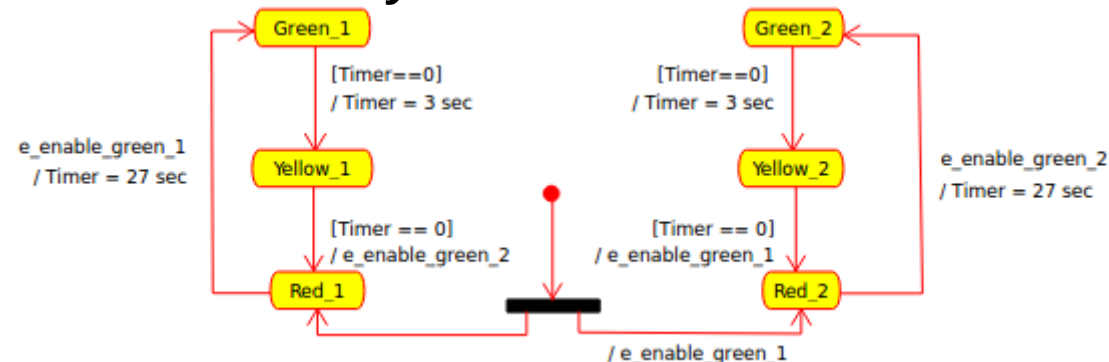

```
//  
// States  
//  
  
class State {  
public:  
    State(int seconds) : _seconds{seconds} { }  
    virtual Traffic_light_color color() {throw runtime_error("State color");}  
  
    State* tic() {  
        if (_seconds > 0) --_seconds;  
        return handle();  
    }  
  
protected:  
    virtual State* handle() {throw runtime_error("State handle");}  
    int _seconds = 0;  
};
```





# A Minor Problem

- Each state must be able to (potentially) create instances of every other state
- We can forward reference *pointers* in C++, but not *instances*
- SOLUTION: We'll create a state\_factory function that generates a state on the heap and returns a pointer, based on a *string* descriptor, e.g., "Green\_1"
  - We'll need to remember to delete each state as we transition away from it to avoid memory leaks



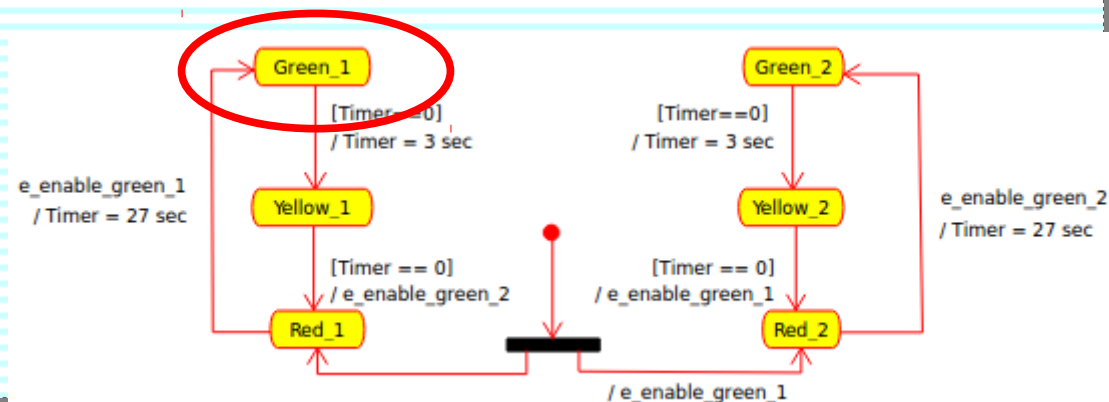

# Our First State - Green\_1

```
// Provide every state access to every other state
State* state_factory(string state);
```

```
class Green_1 : public State {
public:
    Green_1() : State(27) { }
    Traffic_light_color color() override {
        return Traffic_light_color::GREEN;
    }
protected:
    State* handle() override {
        if (_seconds <= 0) {
            return state_factory("Yellow_1");
        } else {
            return this;
        }
    }
};
```

The Green\_1 to Yellow\_1 transition depends only on the timer, not events

Our state\_factory in action







# Yellow\_1 and Red\_1

```
class Yellow_1 : public State {
public:
    Yellow_1() : State(3) { }
    Traffic_light_color color() override {
        return Traffic_light_color::YELLOW;
    }
protected:
    State* handle() override {
        if (_seconds <= 0) {
            e_enable_green_2.generate();
            return state_factory("Red_1");
        } else {
            return this;
        }
    }
};
```

The Yellow\_1 to Red\_1 transition depends only on the timer, but generates an event

```
class Red_1 : public State {
public:
    Red_1() : State(0) { }
    Traffic_light_color color() override {
        return Traffic_light_color::RED;
    }
protected:
    State* handle() override {
        if (e_enable_green_1.consume()) {
            return state_factory("Green_1");
        } else {
            return this;
        }
    }
};
```

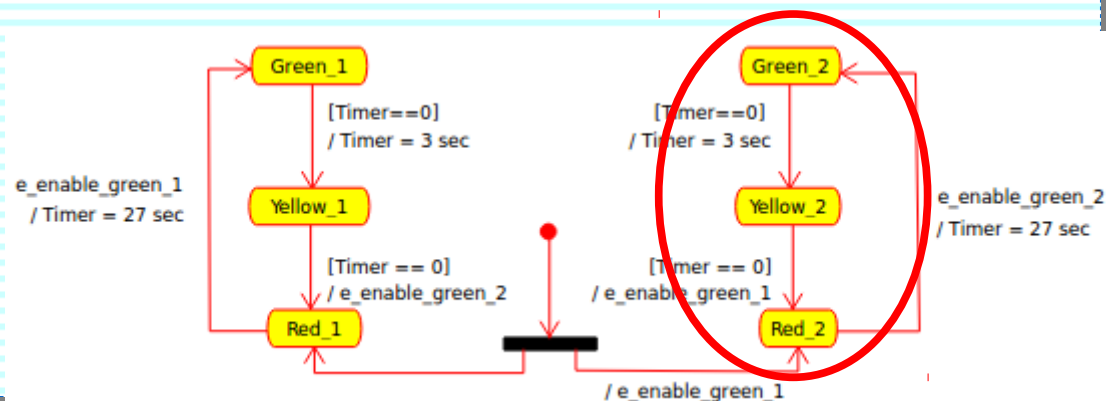

The Red\_1 to Green\_1 transition depends only on an event



# Our state\_factory

```
State* state_factory(string state) {  
    if (state == "Green_1") return new Green_1{};  
    if (state == "Green_2") return new Green_2{};  
    if (state == "Yellow_1") return new Yellow_1{};  
    if (state == "Yellow_2") return new Yellow_2{};  
    if (state == "Red_1") return new Red_1{};  
    if (state == "Red_2") return new Red_2{};  
    throw runtime_error("state_factory: Invalid state: " + state);  
}
```

Green\_2, Yellow\_2, and Red\_2 are very similar – just swap the 1's and 2's. :-)



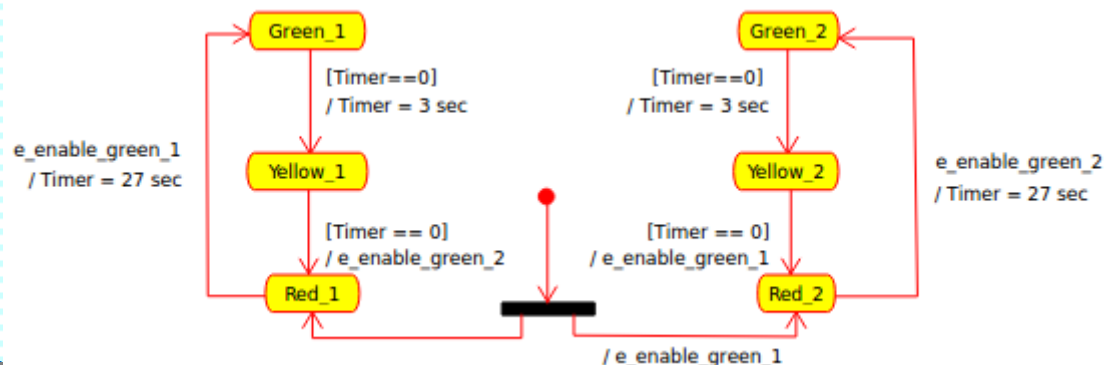



# The State Machine Class

```
//  
// State machines  
//  
  
class Light { // Context  
public:  
    Light(State* state) : _state{state} { }  
    Traffic_light_color color() {return _state->color();}  
    void tic() {  
        State* _newstate = _state->tic();  
        if (_newstate != _state) {  
            delete _state;  
            _state = _newstate;  
        }  
    }  
private:  
    State* _state;  
};
```

The transition logic has been delegated to the states, so the state machine itself is very simple and reusable!

When a state is no longer needed, we must delete it to avoid memory leaks.

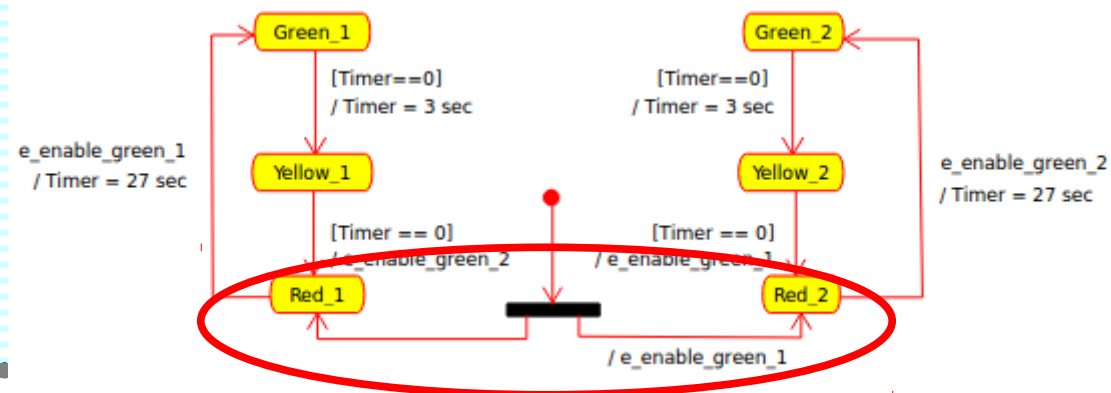



# Finally, main!

```
//////////  
// Main //  
//////////  
int main() {  
    Light north_south{state_factory("Red_1")};  
    Light east_west{state_factory("Red_2")};  
    e_enable_green_1.generate();  
  
    for (int i=0; i < 300; ++i) {  
        cout << i << ": " << ctos(north_south.color()) << " "  
                << ctos(east_west.color()) << endl;  
        east_west.tic();  
        north_south.tic();  
    }  
}
```

Two traffic lights are created,  
distinguished by their initial state.

Then we generate an event to kick  
things off, then tic off 300 seconds.





# Testing

```

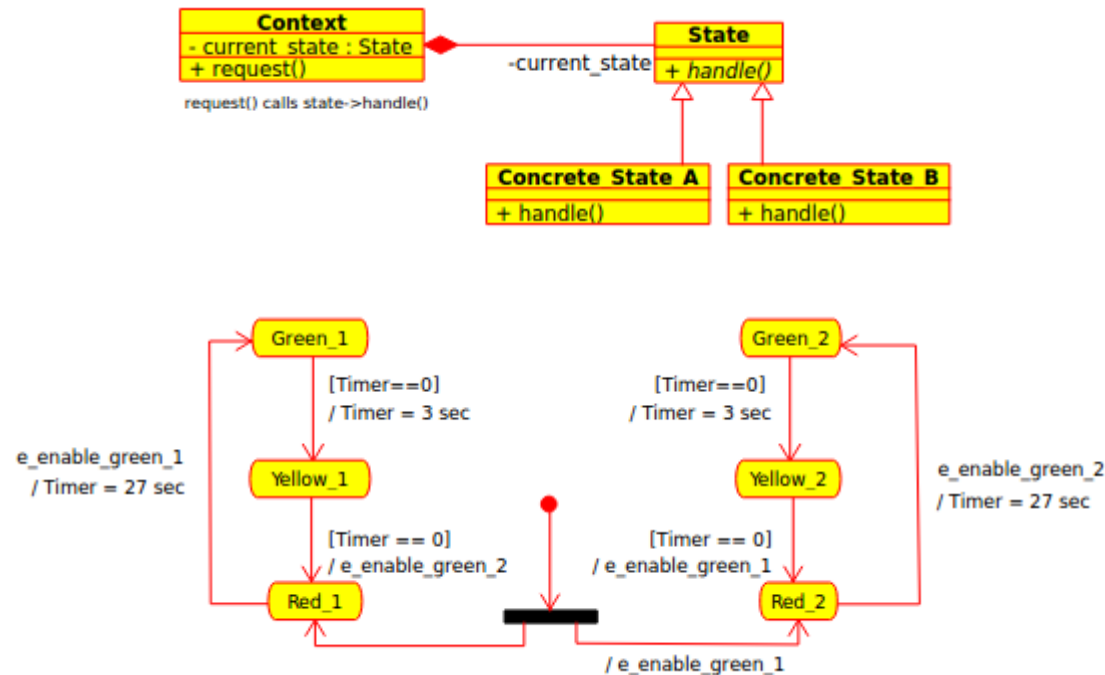
0: red red
1: green red
2: green red
...
25: green red
26: green red
27: green red
28: yellow red
29: yellow red
30: yellow red
31: red red
32: red green
33: red green
34: red green
...
54: red green
55: red green
56: red green
57: red green
58: red green
59: red yellow
60: red yellow
61: red yellow
62: green red
63: green red
64: green red
...
86: green red
87: green red
88: green red
89: yellow red
90: yellow red
91: yellow red
92: red red
93: red green
94: red green
95: red green
...

```

```

117: red green
118: red green
119: red green
120: red yellow
121: red yellow
122: red yellow
123: green red
124: green red
125: green red
...
147: green red
148: green red
149: green red
150: yellow red
151: yellow red
152: yellow red
153: red red
154: red green
155: red green
156: red green
...
178: red green
179: red green
180: red green
181: red yellow
182: red yellow
183: red yellow
184: green red
185: green red
186: green red
...
208: green red
209: green red
210: green red
211: yellow red
212: yellow red
213: yellow red
214: red red
215: red green
216: red green
217: red green
218: red green

```



Output has been truncated at the ellipses to fit multiple cycles on the screen.

```

github.com/AlDanial/cloc v 1.71 T=0.03 s (35.9 files/s, 7141.0 lines/s)
-----
Language          files          blank          comment          code
-----
C++                 1              22             16             161
-----

```



# Other C++ State Machine Implementations

- Professional UML tools offer sophisticated state machine implementations with code generation
  - IBM Rhapsody, MagicDraw, Visual Paradigm...
- Some frameworks provide generalized support
  - Boost MSM, Quantum Platform
- Writing a tailored framework based on the State Design Pattern is always an option





# Quick Review

- A \_\_\_\_\_ is the cumulative value of all relevant stored information to which a system or subsystem has access.
- A \_\_\_\_\_ documents the states, permissible transitions, and activities of a system.
- An \_\_\_\_\_ is a type of occurrence that potentially affects the state of the system
- A \_\_\_\_\_ is a Boolean expression that enables a state transition when true and disables it when false, e.g., [power == on]
- True or False: A UML State Diagram is a special case of the UML Activity Diagram.
- The \_\_\_\_\_ supports full encapsulation of unlimited states within a scalable context. It does not, however, support events as is. It is closely related to the \_\_\_\_\_.
- A \_\_\_\_\_ state machine defines outputs based on current state and current inputs, while a \_\_\_\_\_ state machine defines outputs based on current state only.



# For Next Class

- (Optional) Read Chapter 25 in Stroustrup
  - Do the Drills!
- Skim Chapters 17-19 for next week
- **Sprint #3 now in progress: Our first ice cream order!**
  - **Individuals:** Time to create your servers and customers, and package it all together into your first multi-serving order!
  - **Duos:** Also create the Emporium itself (perhaps via an emporium class), and start loading and saving data files!
  - **Trios:** Save and load data files, and add some enticing photos; and the manager looks forward to the first reports!
  - **Quattros:** *Lots* of reports, enticing photos, restocking the serving prep counter, and then pay those servers!